

Microservices: Decomposing Applications for Deployability and Scalability

Chris Richardson

Author of POJOs in Action

Founder of the original CloudFoundry.com

 @crichtson

chris@chrisrichardson.net

<http://plainoldobjects.com>



The Java Event
for ALL Developers

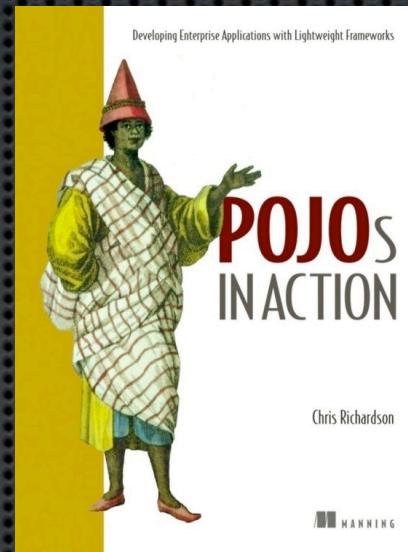
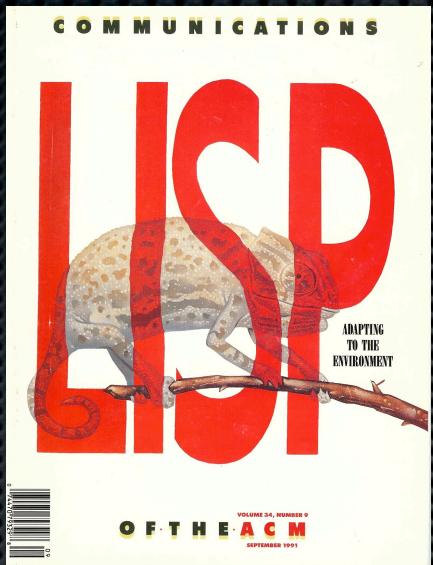
Presentation goal

How decomposing applications
improves deployability and
scalability

and

simplifies the adoption of new
technologies

About Chris

A screenshot of the Cloud Foundry website. The header includes fields for 'Email' and 'Password', and links for 'Sign Up', 'Forgot password?', and 'SIGN IN'. Below the header, there are navigation links for 'HOW WE HELP', 'FEATURES', 'INFORMATION', 'BLOG', and 'CONTACT US'. A system alert message states: 'SYSTEM ALERT. PLEASE READ: Cloud Foundry will be moving to a new URL. [More](#)' with a small red exclamation mark icon. The main content area features a green gradient background. On the left, it says 'The Enterprise Java Cloud' and lists three bullet points: 'Real Java Applications Deployed in Minutes', 'Built for Spring and Grails Web Applications', and 'Most Widely Used Technologies Delivered as a Platform'. It includes 'SIGN UP' and 'LEARN MORE' buttons. In the center, there is a black box with the 'CLOUDFOUNDRY SPRINGSOURCE' logo and a play button icon, with the text 'APPLICATION DEMO Deploying Web Applications To Amazon EC2 with Cloud Foundry' below it.

@crichtudson

About Chris

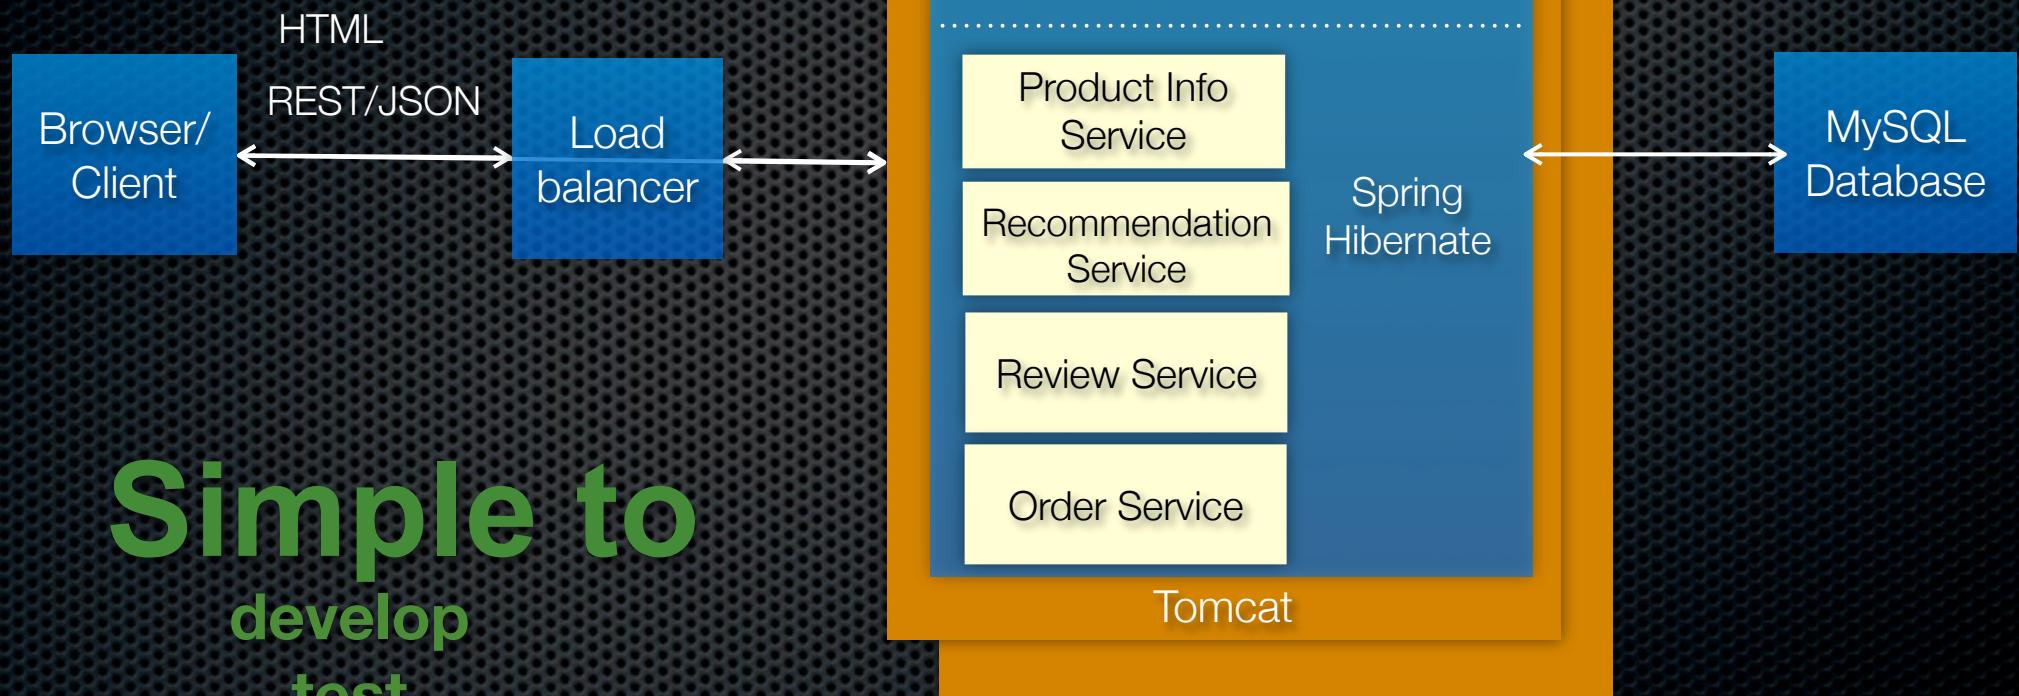
- ❖ Founder of a buzzword compliant (stealthy, social, mobile, big data, machine learning, ...) startup
- ❖ Consultant helping organizations improve how they architect and deploy applications using cloud computing, micro services, polyglot applications, NoSQL, ...

Agenda

- The (sometimes evil) monolith
- Decomposing applications into services
- Client ⇔ service interaction design
- Decentralized data management

Let's imagine you are
building an online store

Traditional application architecture



Simple to
develop
test
deploy
scale

But large, complex, monolithic
applications



problems

Intimidates developers



Obstacle to frequent deployments

- Need to redeploy everything to change one component
- Interrupts long running background (e.g. Quartz) jobs
- Increases risk of failure

Eggs in
one basket



- Updates will happen less often - really long QA cycles
- e.g. Makes A/B testing UI really difficult

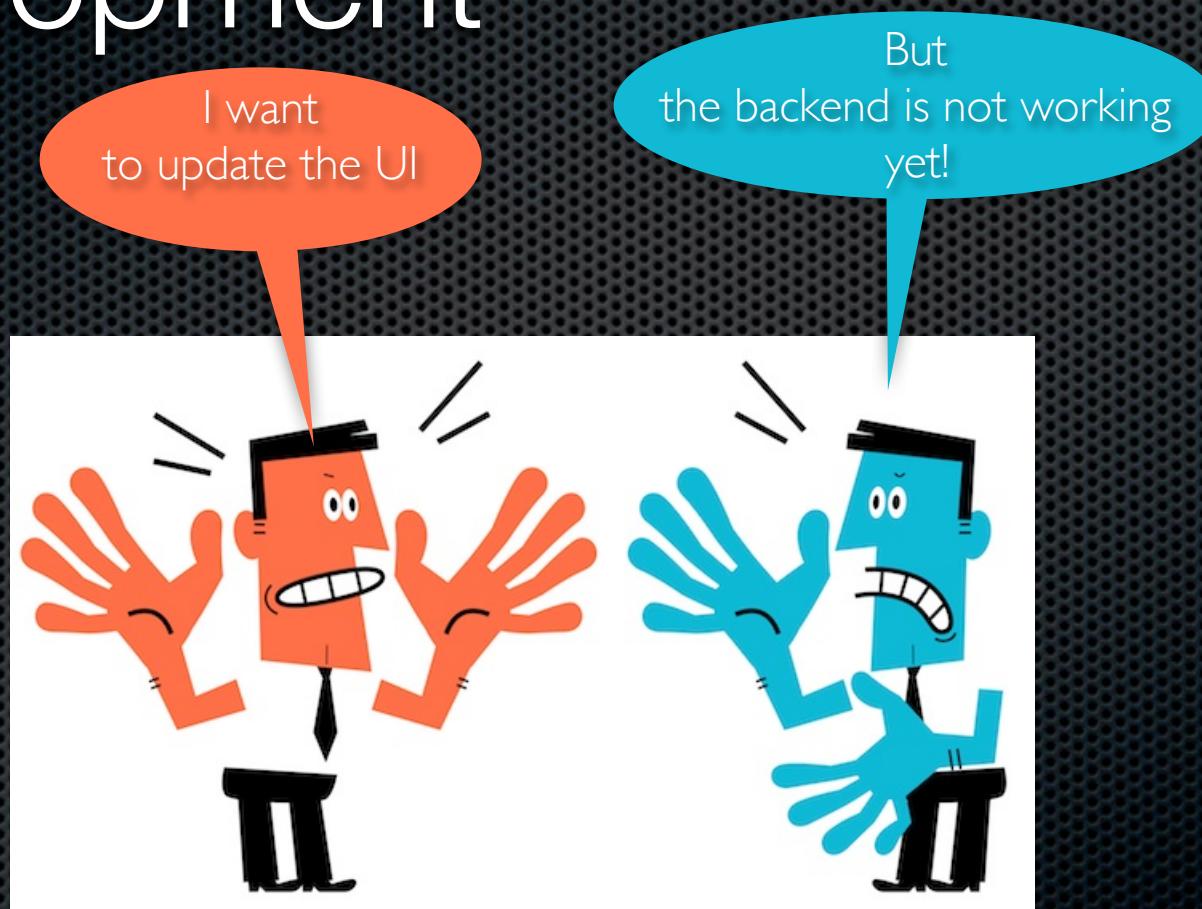
Overloads your IDE and container



Slows down development

@crichtson

Obstacle to scaling development



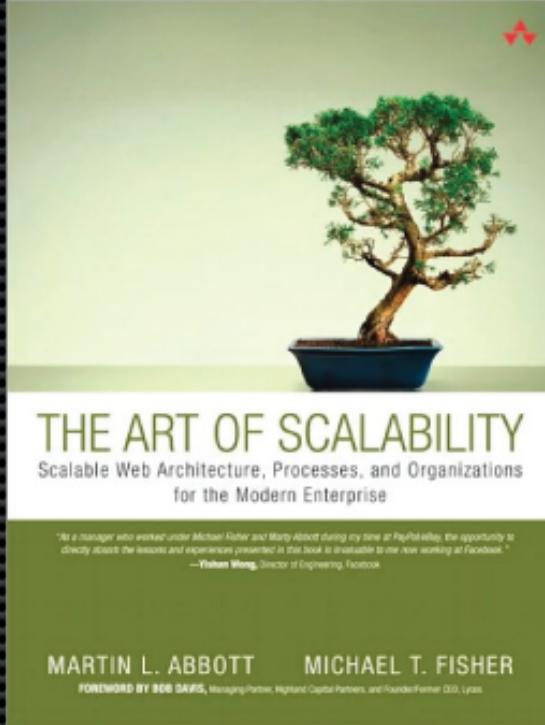
Lots of coordination and communication required

Requires long-term commitment to a technology stack



Agenda

- The (sometimes evil) monolith
- Decomposing applications into services
- Client ⇔ service interaction design
- Decentralized data management

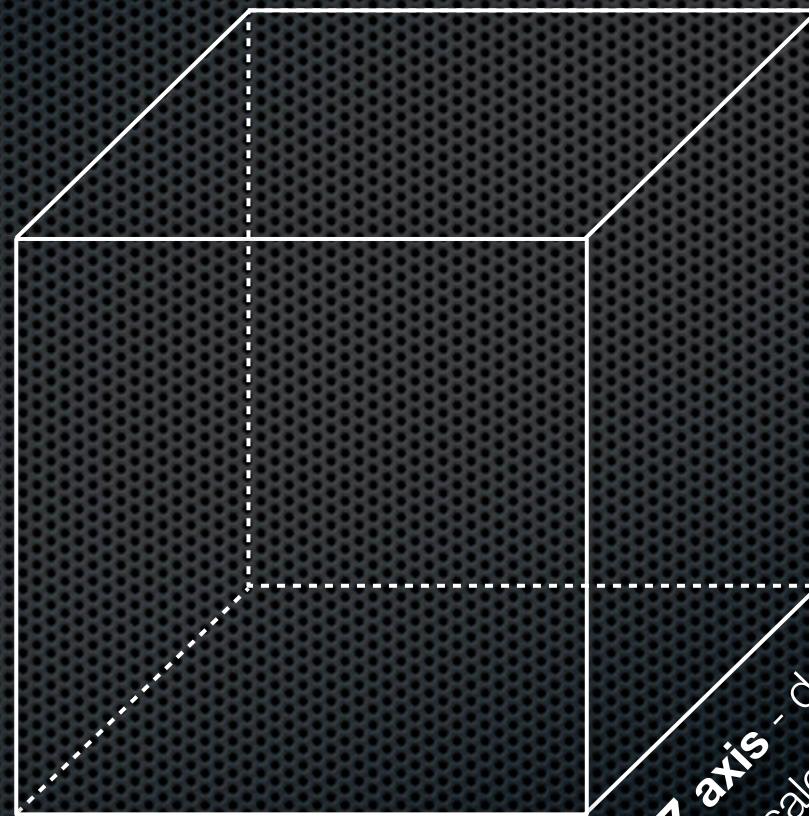


@crichardson

The scale cube

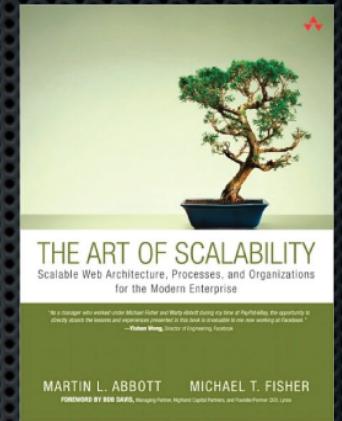
Y axis -
functional
decomposition

Scale by
splitting
different things



X axis
- horizontal duplication

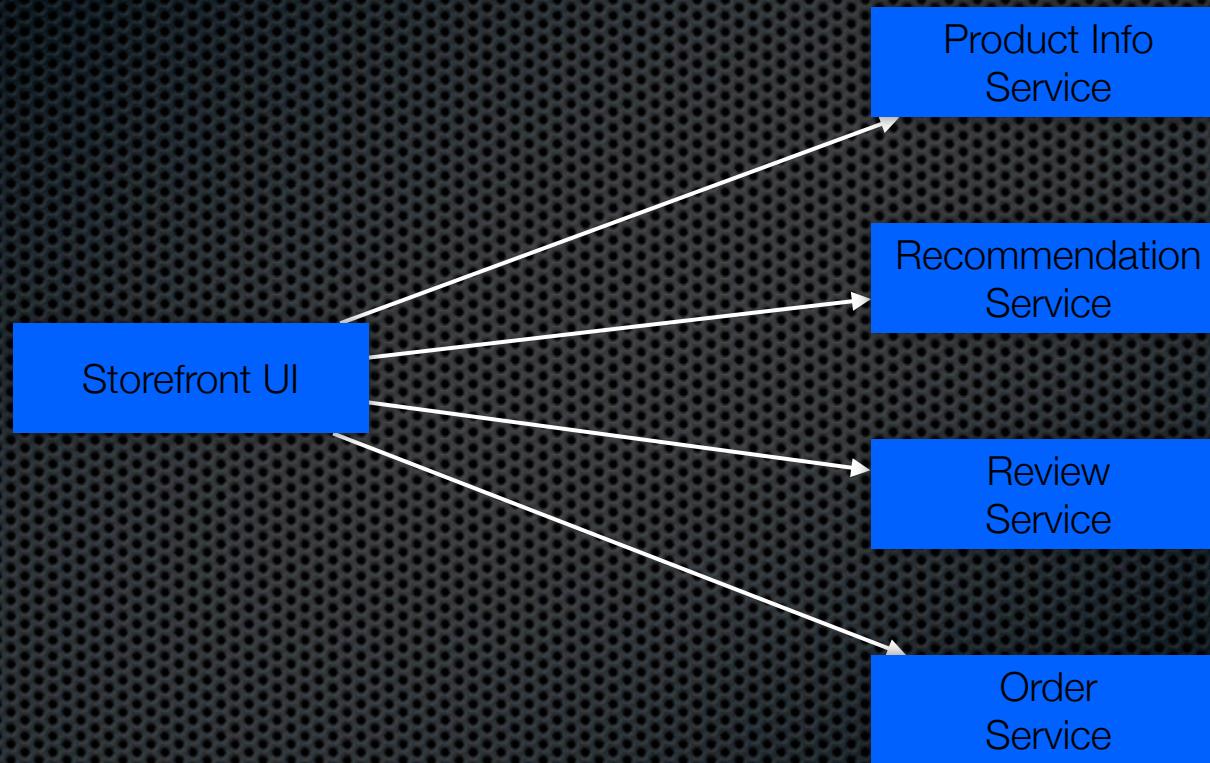
Z axis - data partitioning
Scale by splitting similar
things



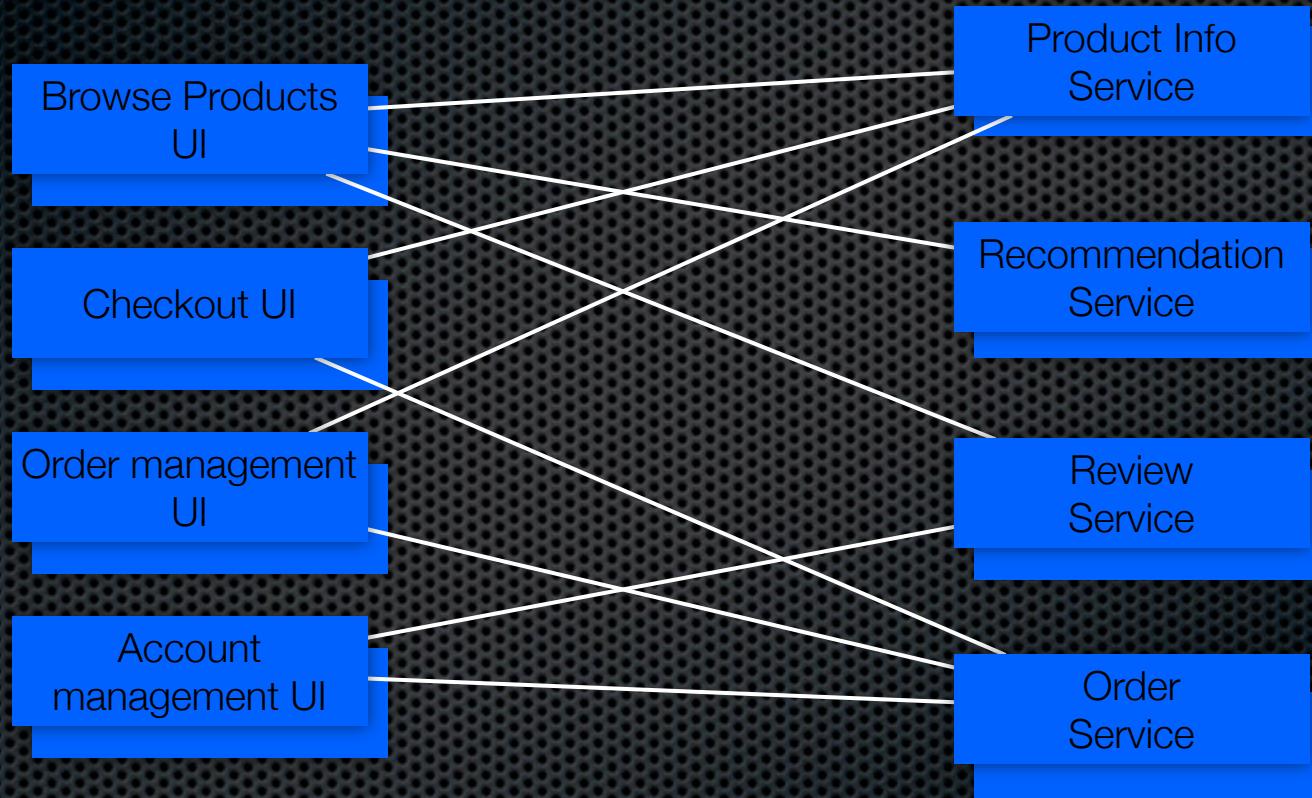
Y-axis scaling - application level



Y-axis scaling - application level



Y-axis scaling - application level



Apply X-axis and Z-axis scaling
to each service independently

Service deployment options

Isolation, manageability

- VM or Physical Machine
- Docker/Linux container
- JVM
- JAR/WAR/OSGI bundle/...



Density/efficiency

Partitioning strategies...

- ❖ Partition by noun, e.g. product info service
- ❖ Partition by verb, e.g. Checkout UI
- ❖ Single Responsibility Principle
- ❖ Unix utilities - do one focussed thing well

Partitioning strategies

- ❖ Too few
 - ❖ Drawbacks of the monolithic architecture
- ❖ Too many - a.k.a. Nano-service anti-pattern
 - ❖ Runtime overhead
 - ❖ **Potential** risk of excessive network hops
 - ❖ **Potentially** difficult to understand system

Something of an art

Example micro-service using Spring Boot



Rob Winch (@rob_winch)

@Controller

```
class ThisWillActuallyRun {  
    @RequestMapping("/")  
    @ResponseBody  
    String home() {  
        "Hello World!"  
    }  
}
```

Reply Retweet Favorite Buffer More

RETWEETS 65 FAVORITES 18

3:12 PM - 6 Aug 2013

A screenshot of a Twitter post from user @rob_winch. The post contains Java code for a Spring Boot controller named 'ThisWillActuallyRun'. The code defines a single endpoint at the root path that returns the string "Hello World!". The tweet has received 65 retweets and 18 favorites. The timestamp is 3:12 PM - 6 Aug 2013.

For more on micro-services see
<http://microservices.io>

@crichtson

But more realistically...

Focus on building cohesive services that make development and deployment easier

- not just tiny services

Real world examples



<http://techblog.netflix.com/>

~600 services



<http://highscalability.com/amazon-architecture>

100-150 services to build a page



<http://www.addsimplicity.com/downloads/eBaySDForum2006-11-29.pdf>

<http://queue.acm.org/detail.cfm?id=1394128>

There are drawbacks

Complexity of developing and managing a distributed system

<http://contino.co.uk/blog/2013/03/31/microservices-no-free-lunch.html>

*Using a PaaS can significantly
simplify deployment*

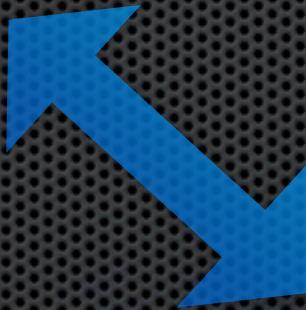
Multiple databases & Transaction management

Deploying features that span
multiple services requires
careful coordination

When to use it?

In the beginning:

- You don't need it
- It will slow you down



Later on:

- You need it
- Refactoring is painful

But there are many benefits

Smaller, simpler apps

- Easier to understand and develop
- Reduced startup time - important for GAE
- Less jar/classpath hell - who needs OSGI?

Scales development:
develop, deploy and scale
each service independently

Improves fault isolation

Eliminates long-term commitment
to a single technology stack



Modular, polyglot, multi-
framework applications

Two levels of architecture

System-level

Services

Inter-service glue: interfaces and communication mechanisms

Slow changing

Service-level

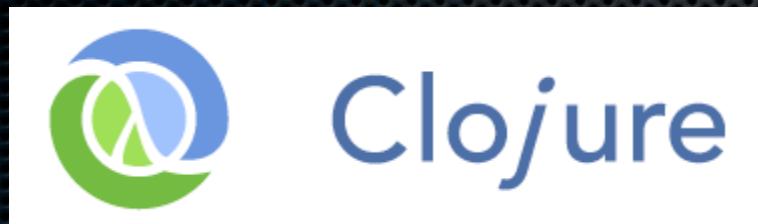
Internal architecture of each service

Each service could use a different technology stack

Pick the best tool for the job

Rapidly evolving

Easily try other technologies



Spring Boot



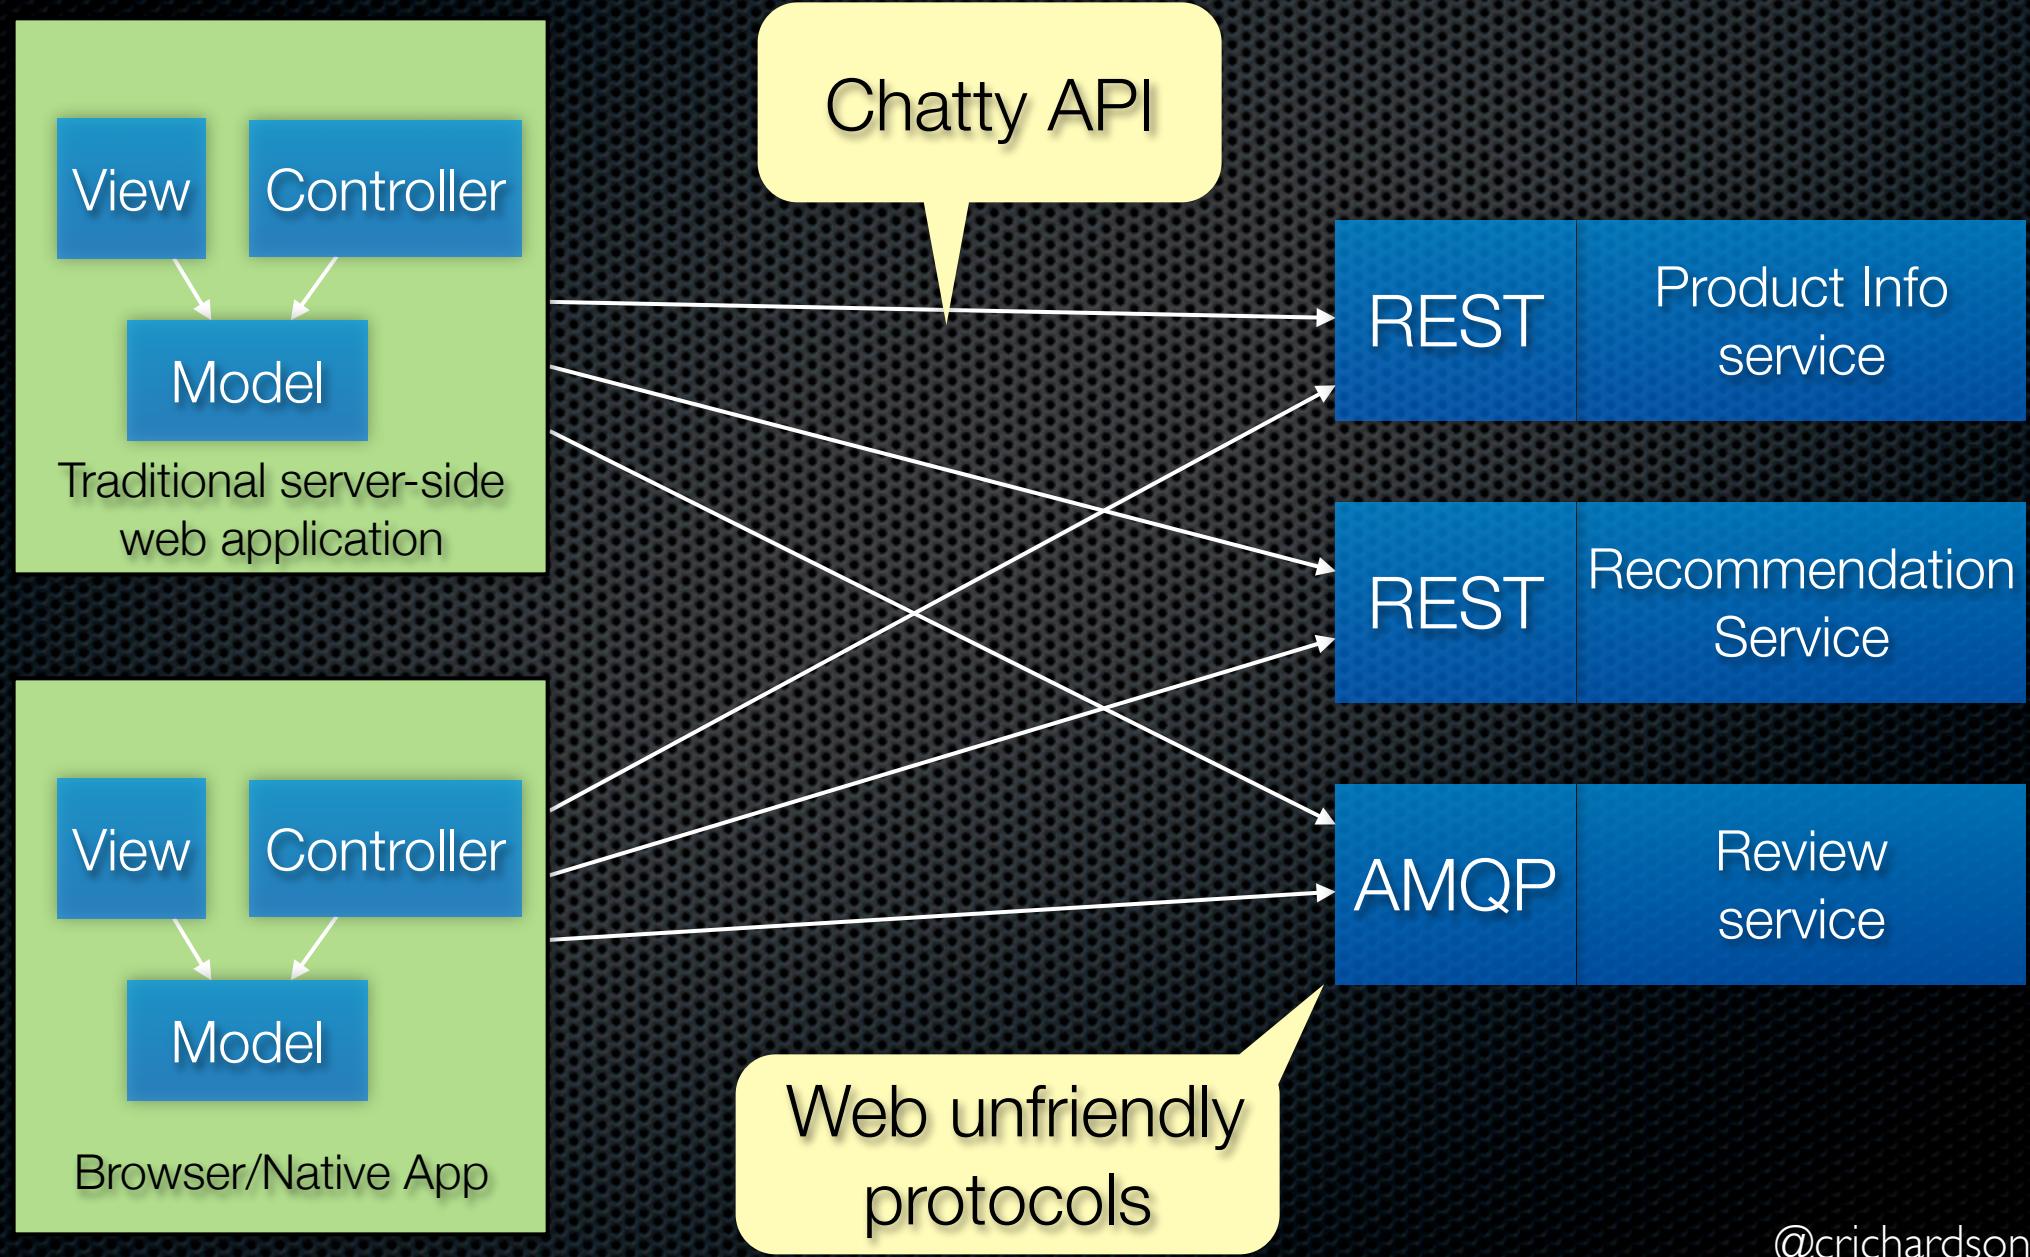
... and fail safely

Agenda

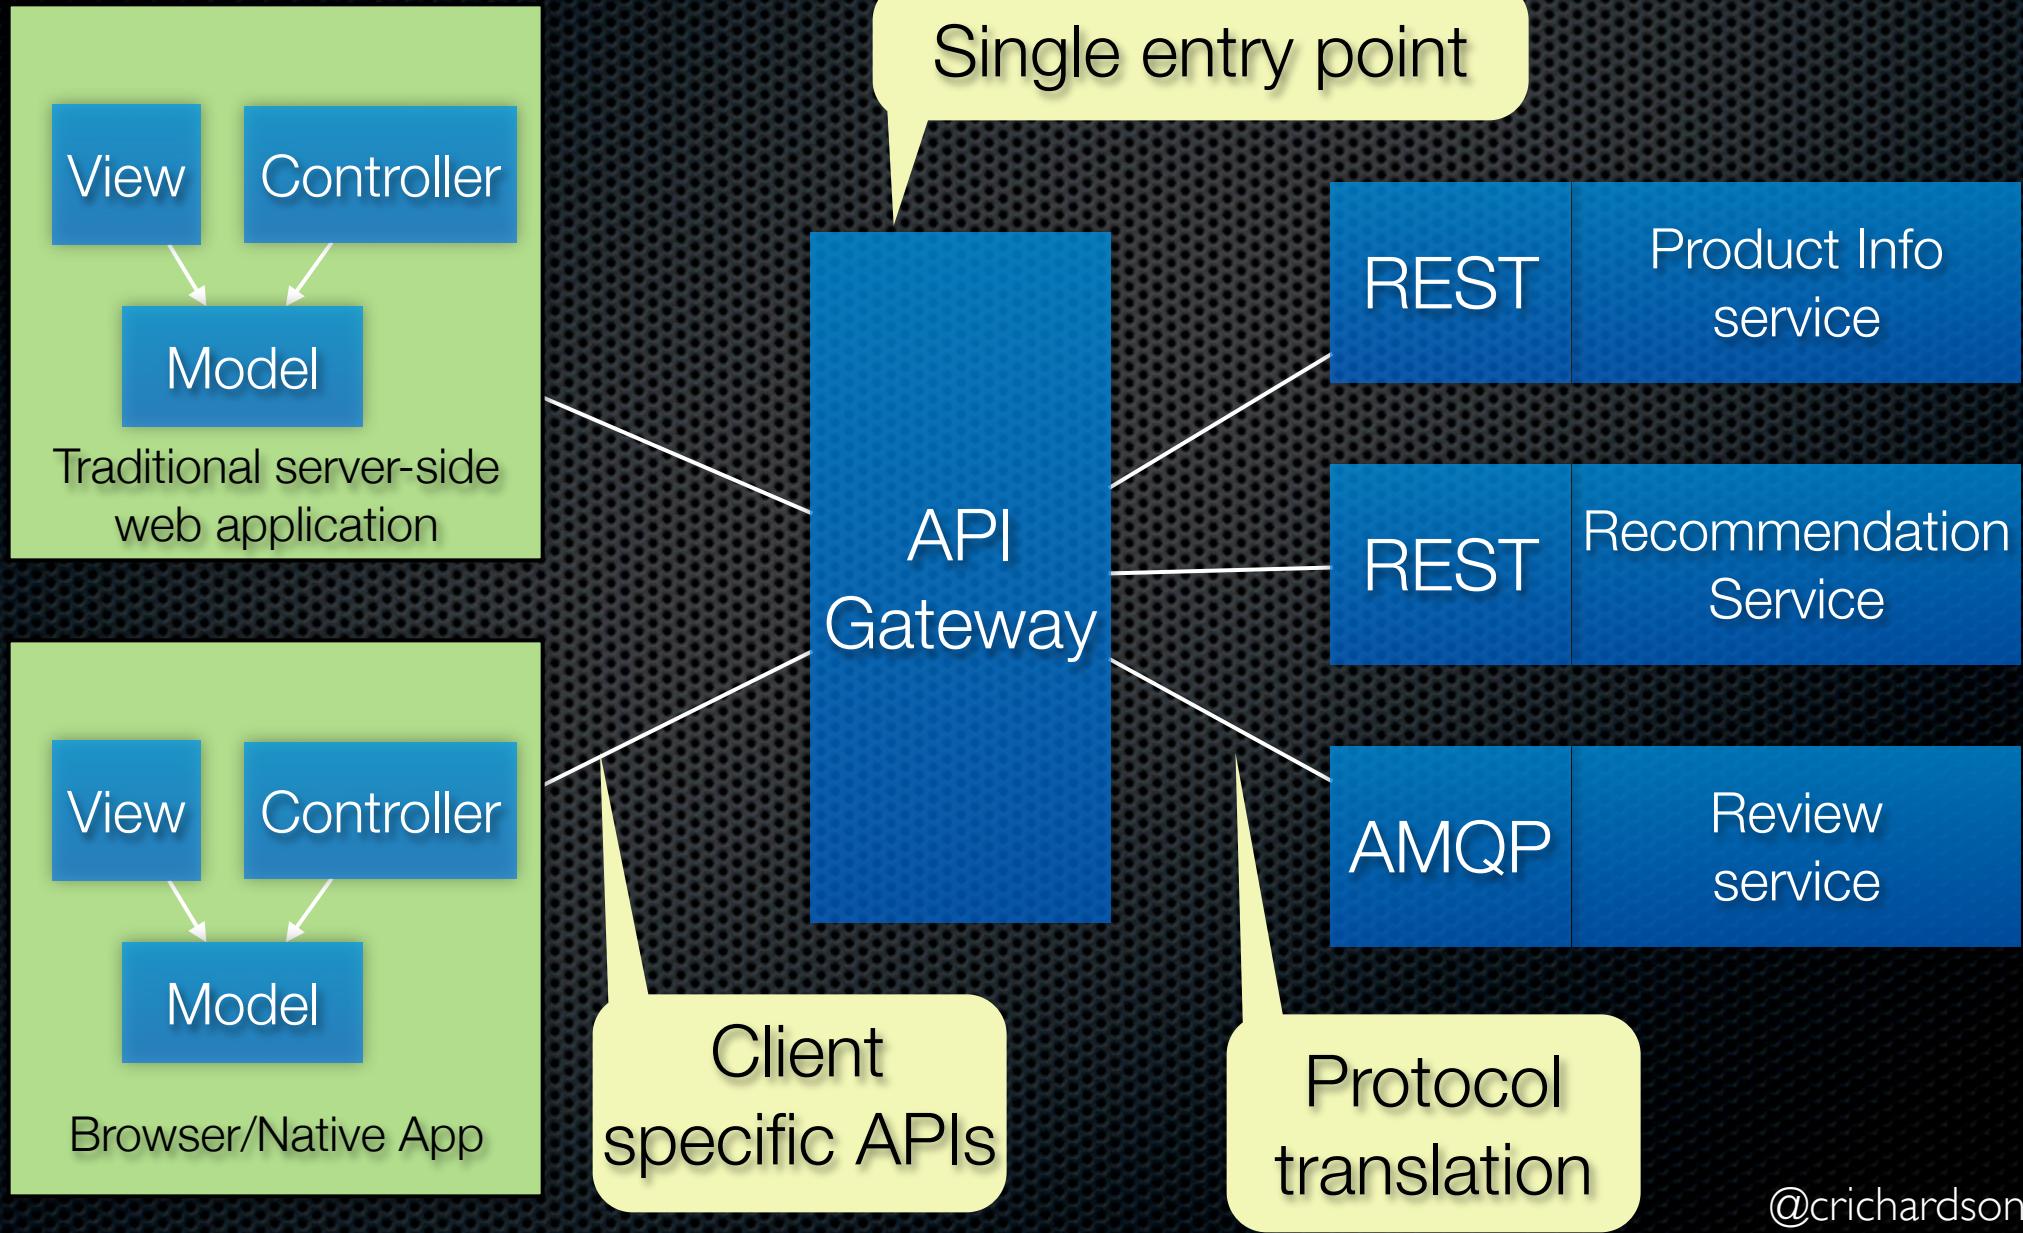
- The (sometimes evil) monolith
- Decomposing applications into services
- Client ⇔ service interaction design
- Decentralized data management

How do clients of the system
interact with the services?

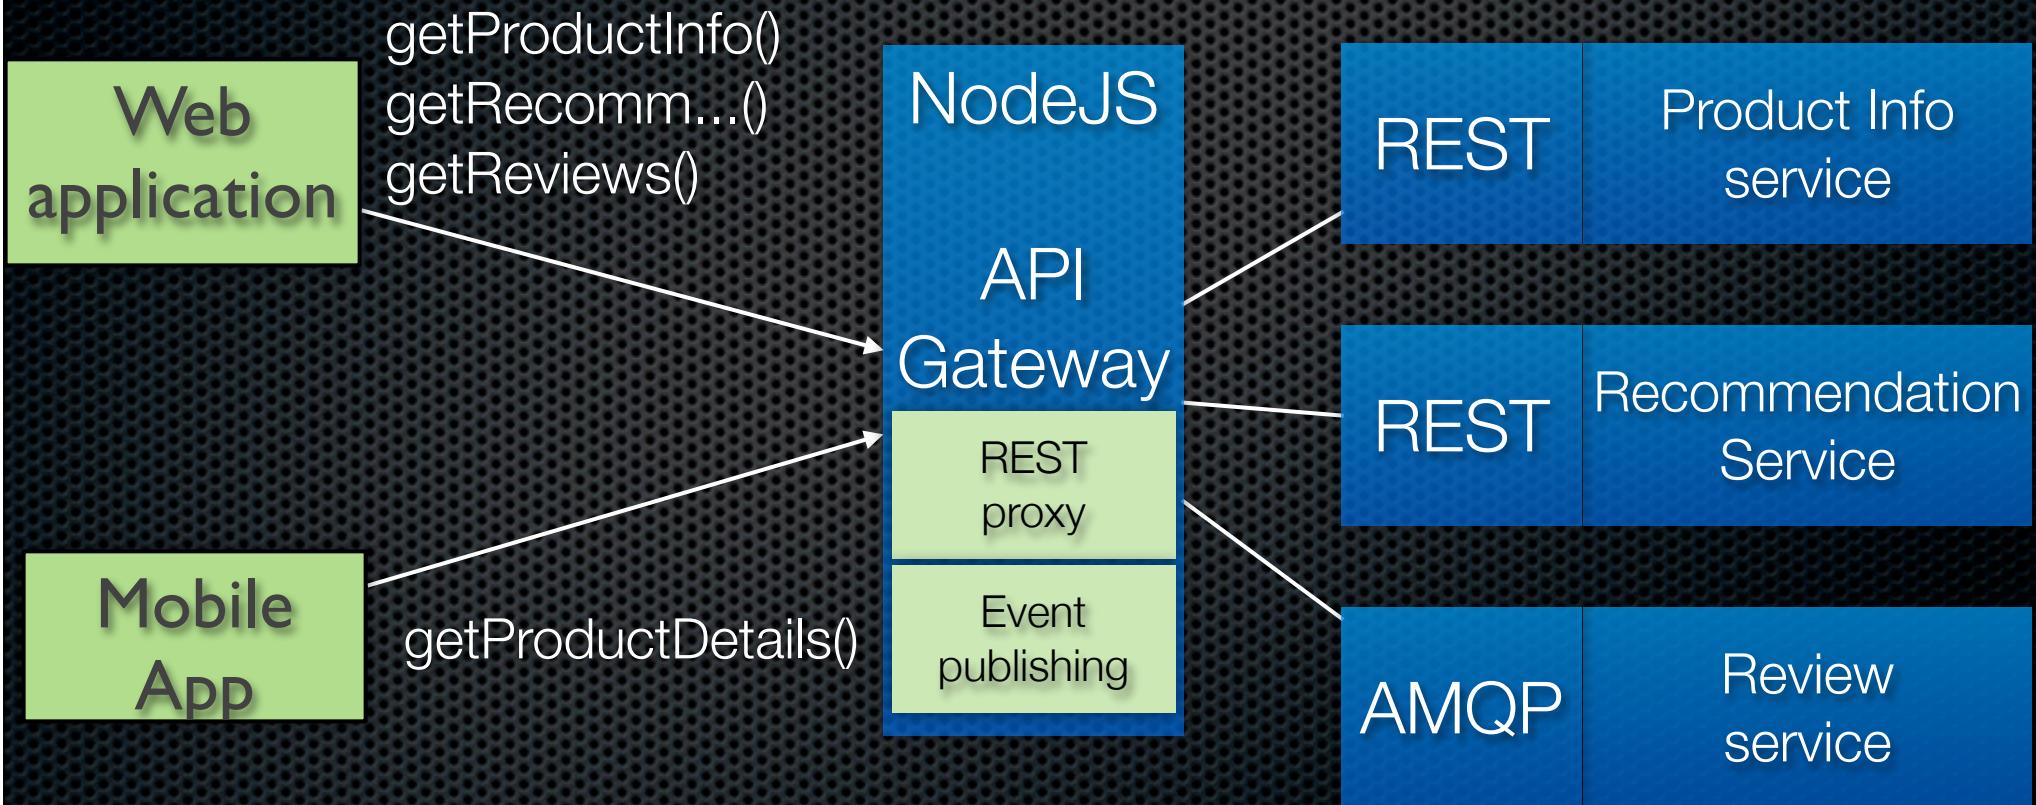
Directly connecting the front-end to the backend



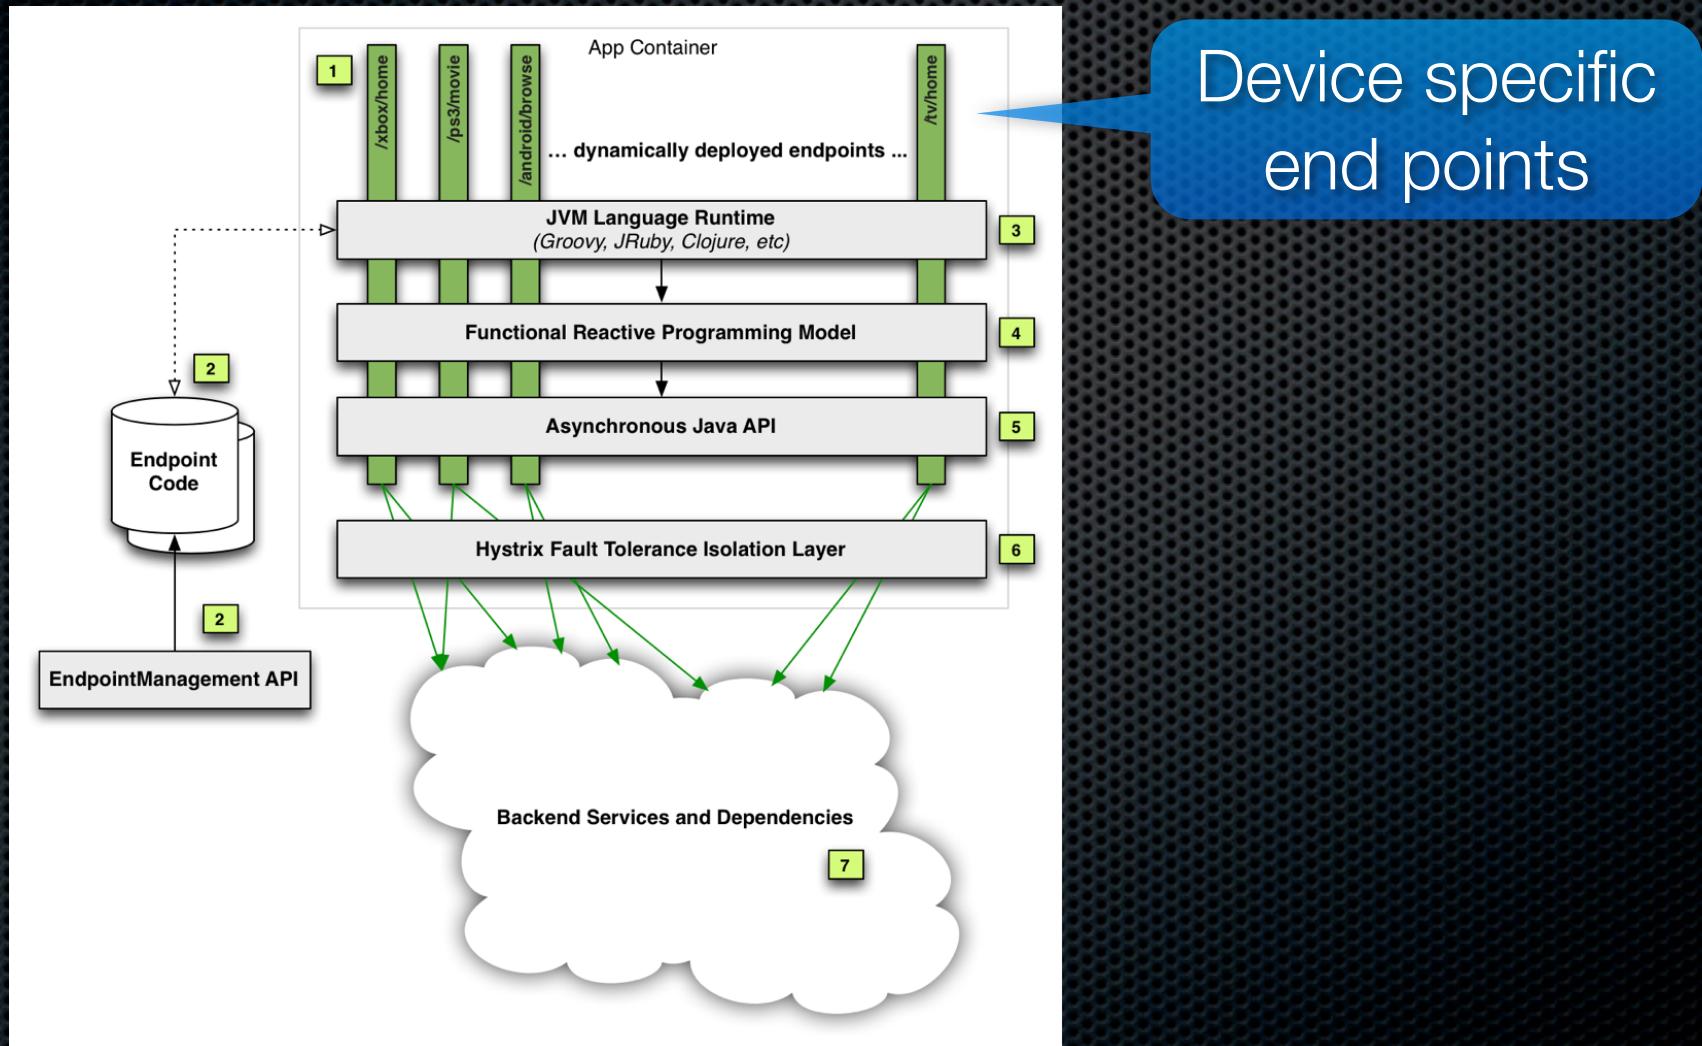
Use an API gateway



Optimized client-specific APIs



Netflix API Gateway



<http://techblog.netflix.com/2013/01/optimizing-netflix-api.html>

@crichton

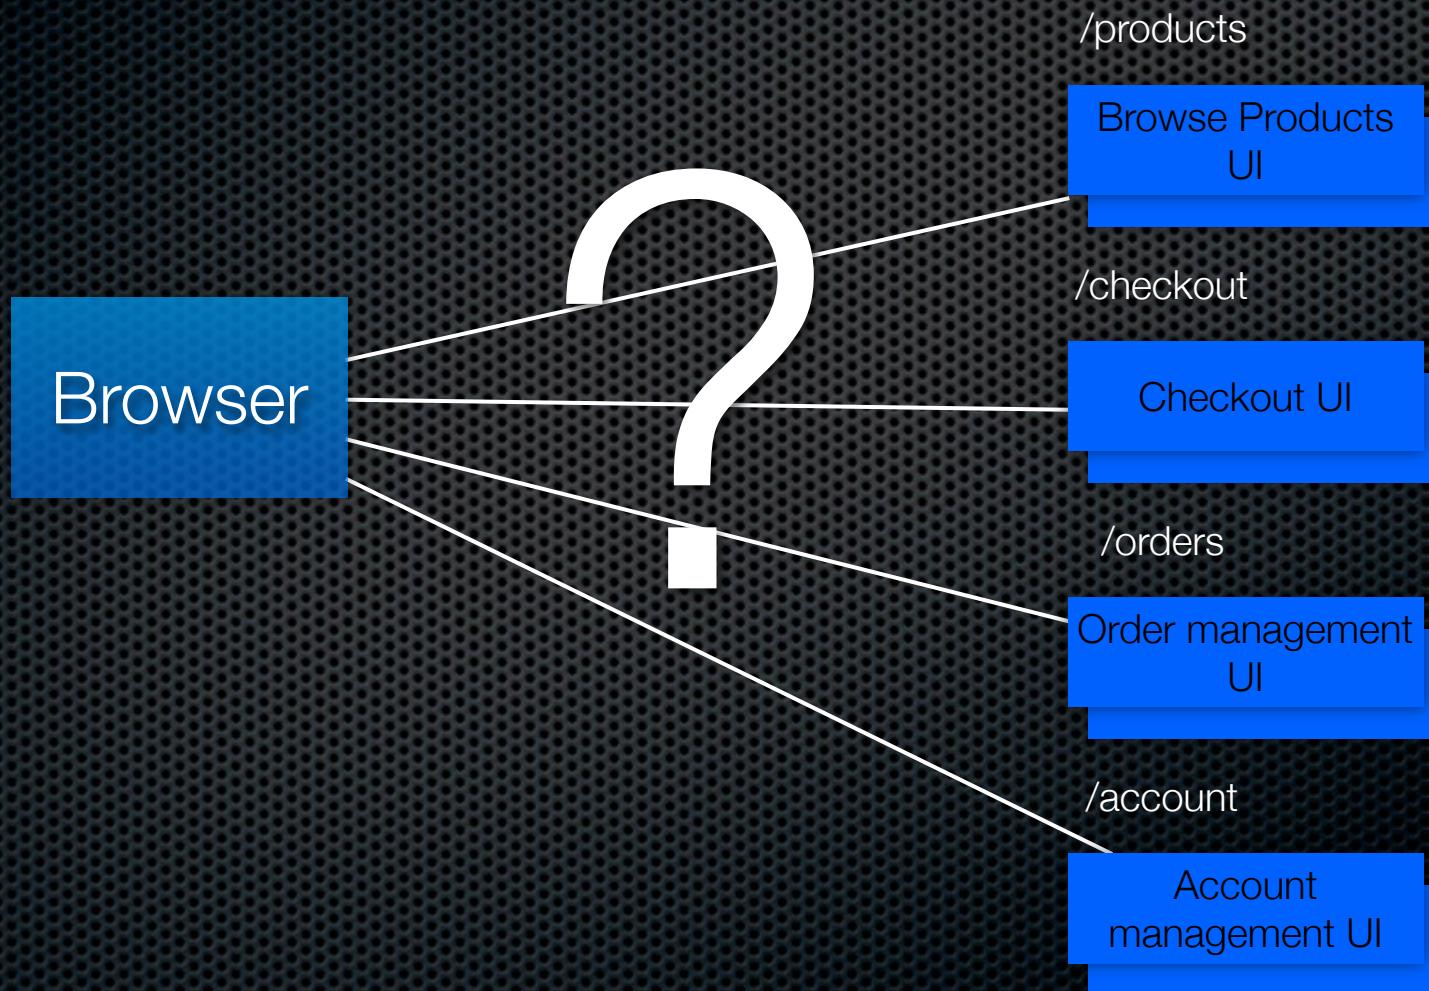
API gateway design challenges

- ❖ Performance and scalability
 - ❖ Non-blocking I/O
 - ❖ Asynchronous, concurrent code
- ❖ Handling partial failures
- ❖

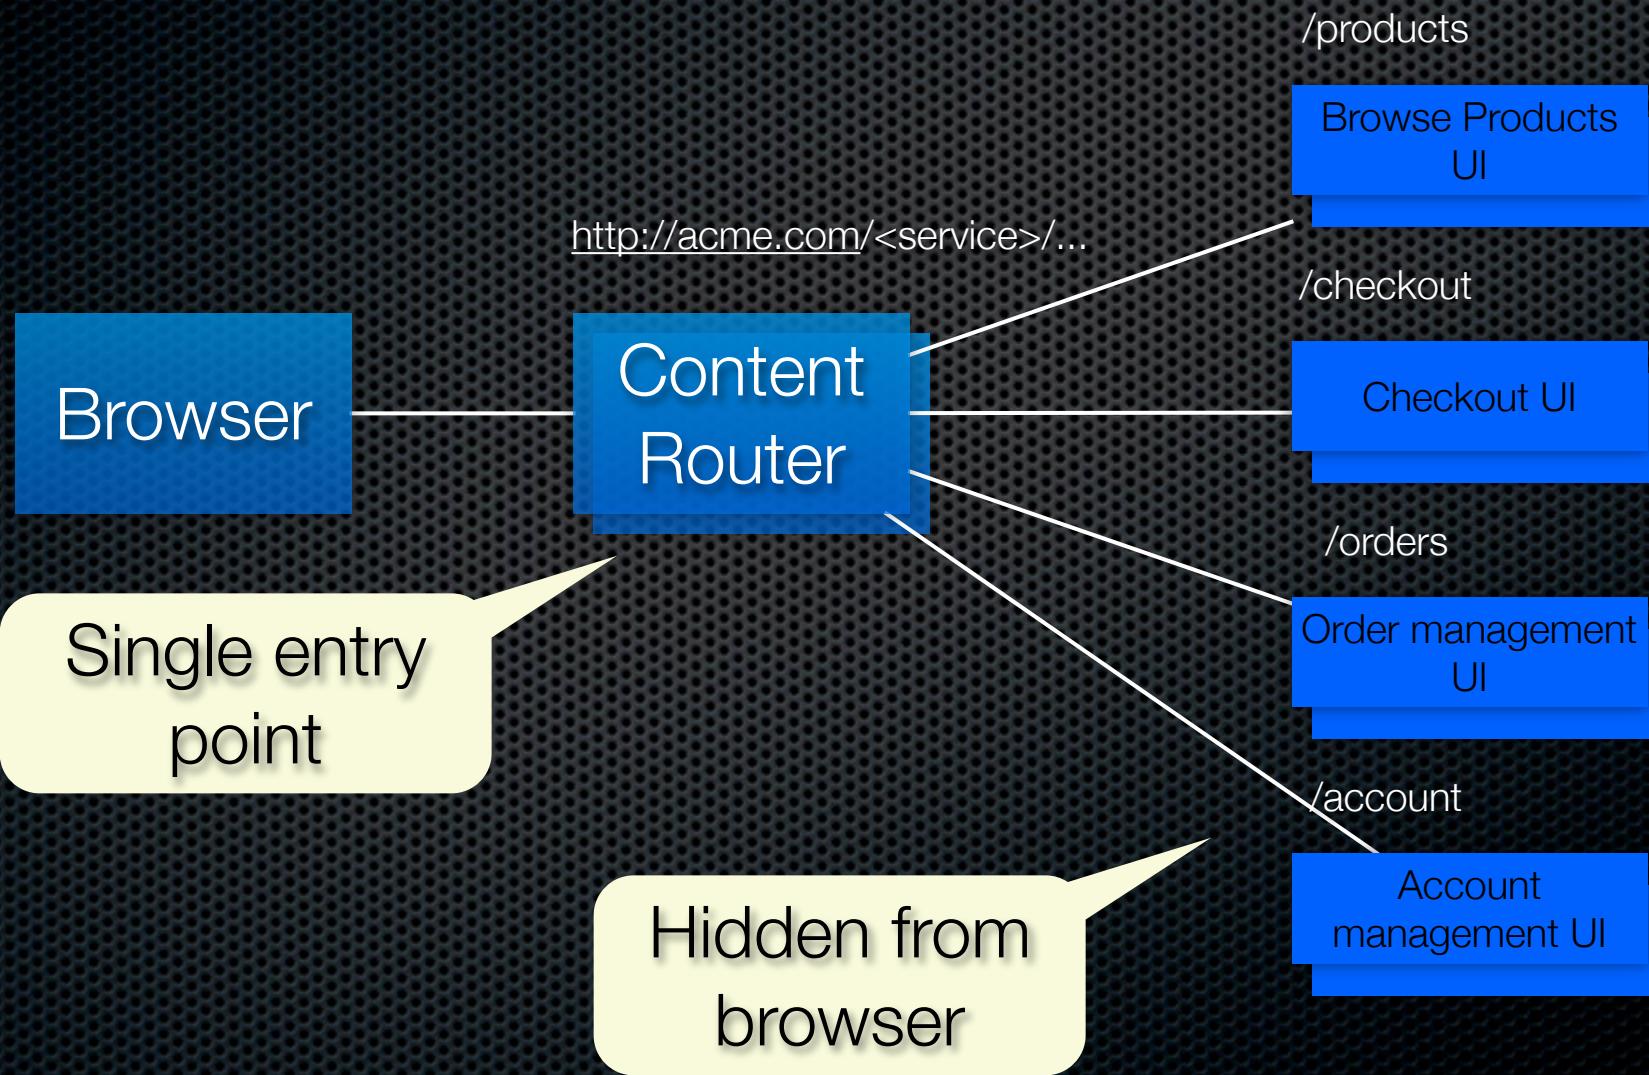
<http://techblog.netflix.com/2012/02/fault-tolerance-in-high-volume.html>

How does a browser
interact with the partitioned
web application?

Partitioned web app ⇒ no longer a single base URL



The solution: single entry point that routes based on URL



How do the services communicate?

Inter-service communication options

- Synchronous HTTP ⇔ asynchronous AMQP
- Formats: JSON, XML, Protocol Buffers, Thrift, ...

Asynchronous is preferred
JSON is fashionable but binary format
is more efficient

Pros and cons of messaging

Pros

- ▣ Decouples client from server
- ▣ Message broker buffers messages
- ▣ Supports a variety of communication patterns

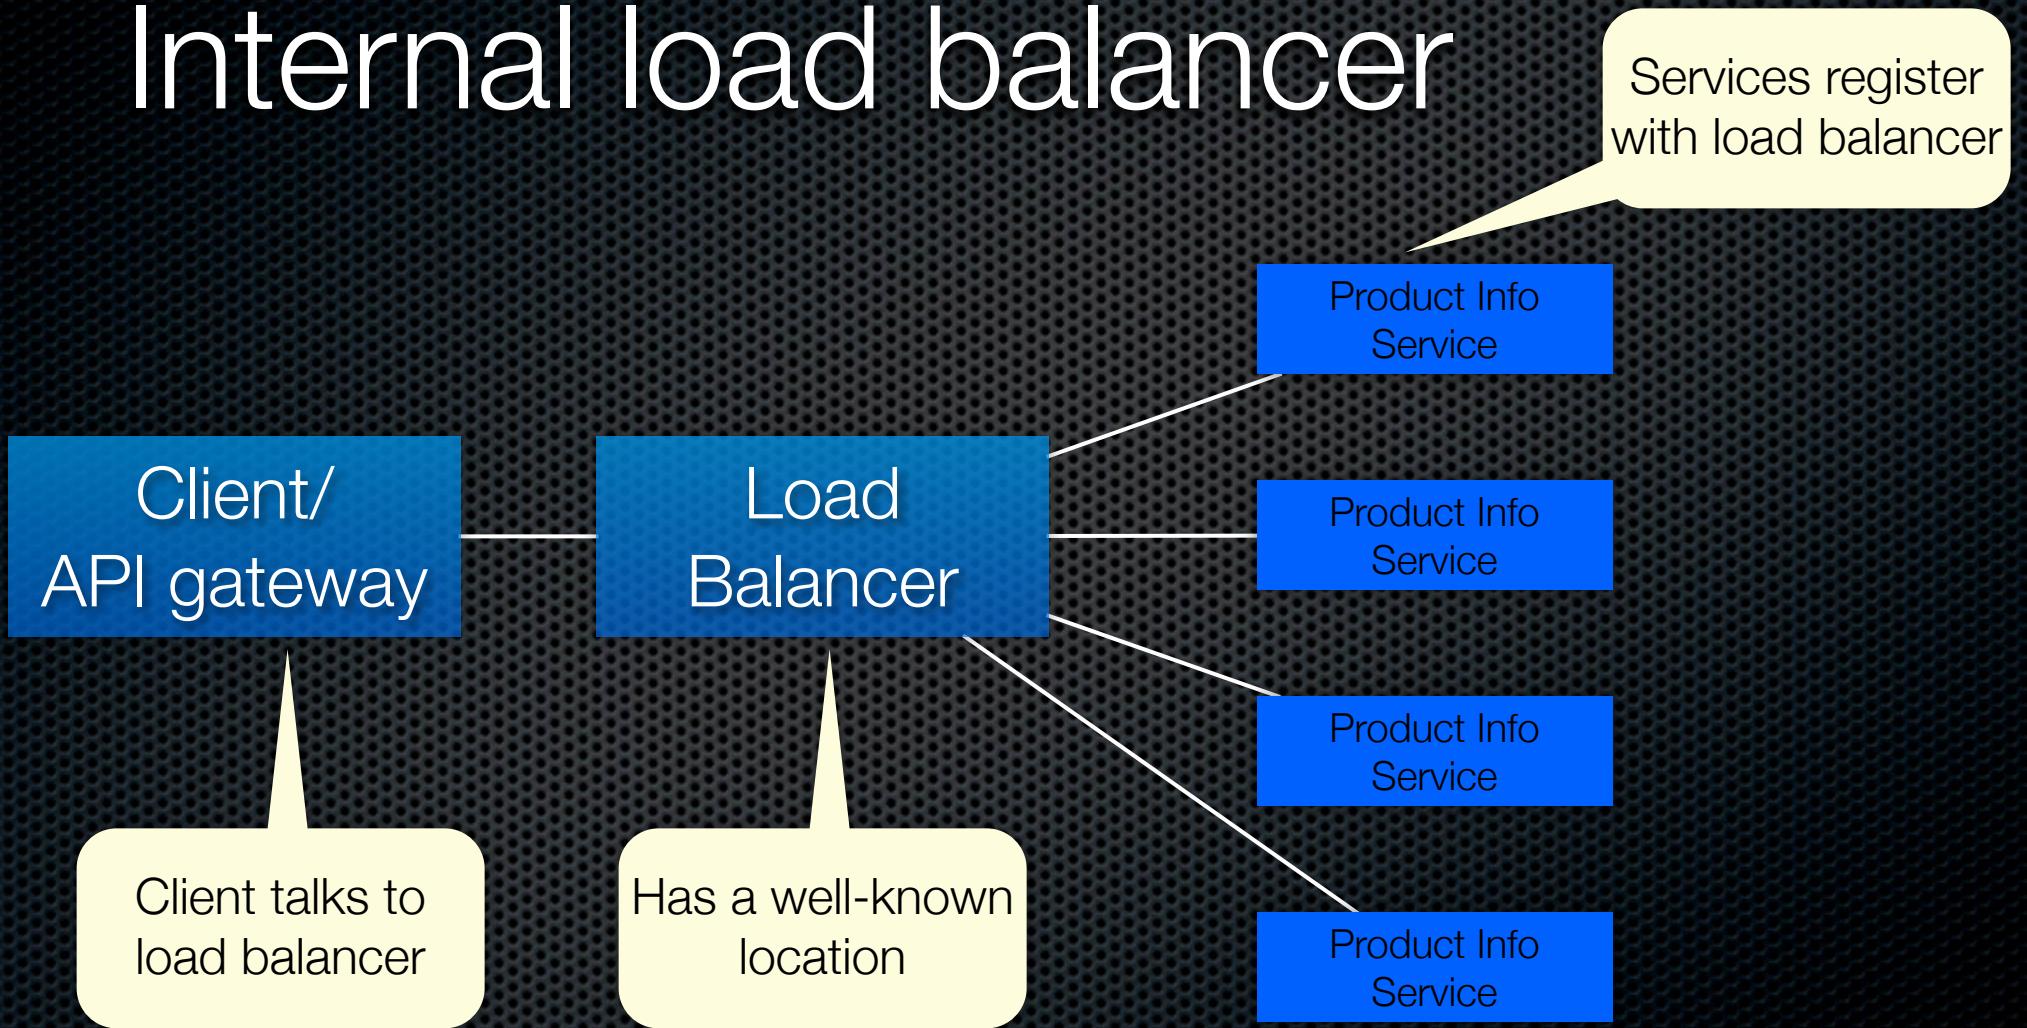
Cons

- ▣ Additional complexity of message broker
- ▣ Request/reply-style communication is more complex

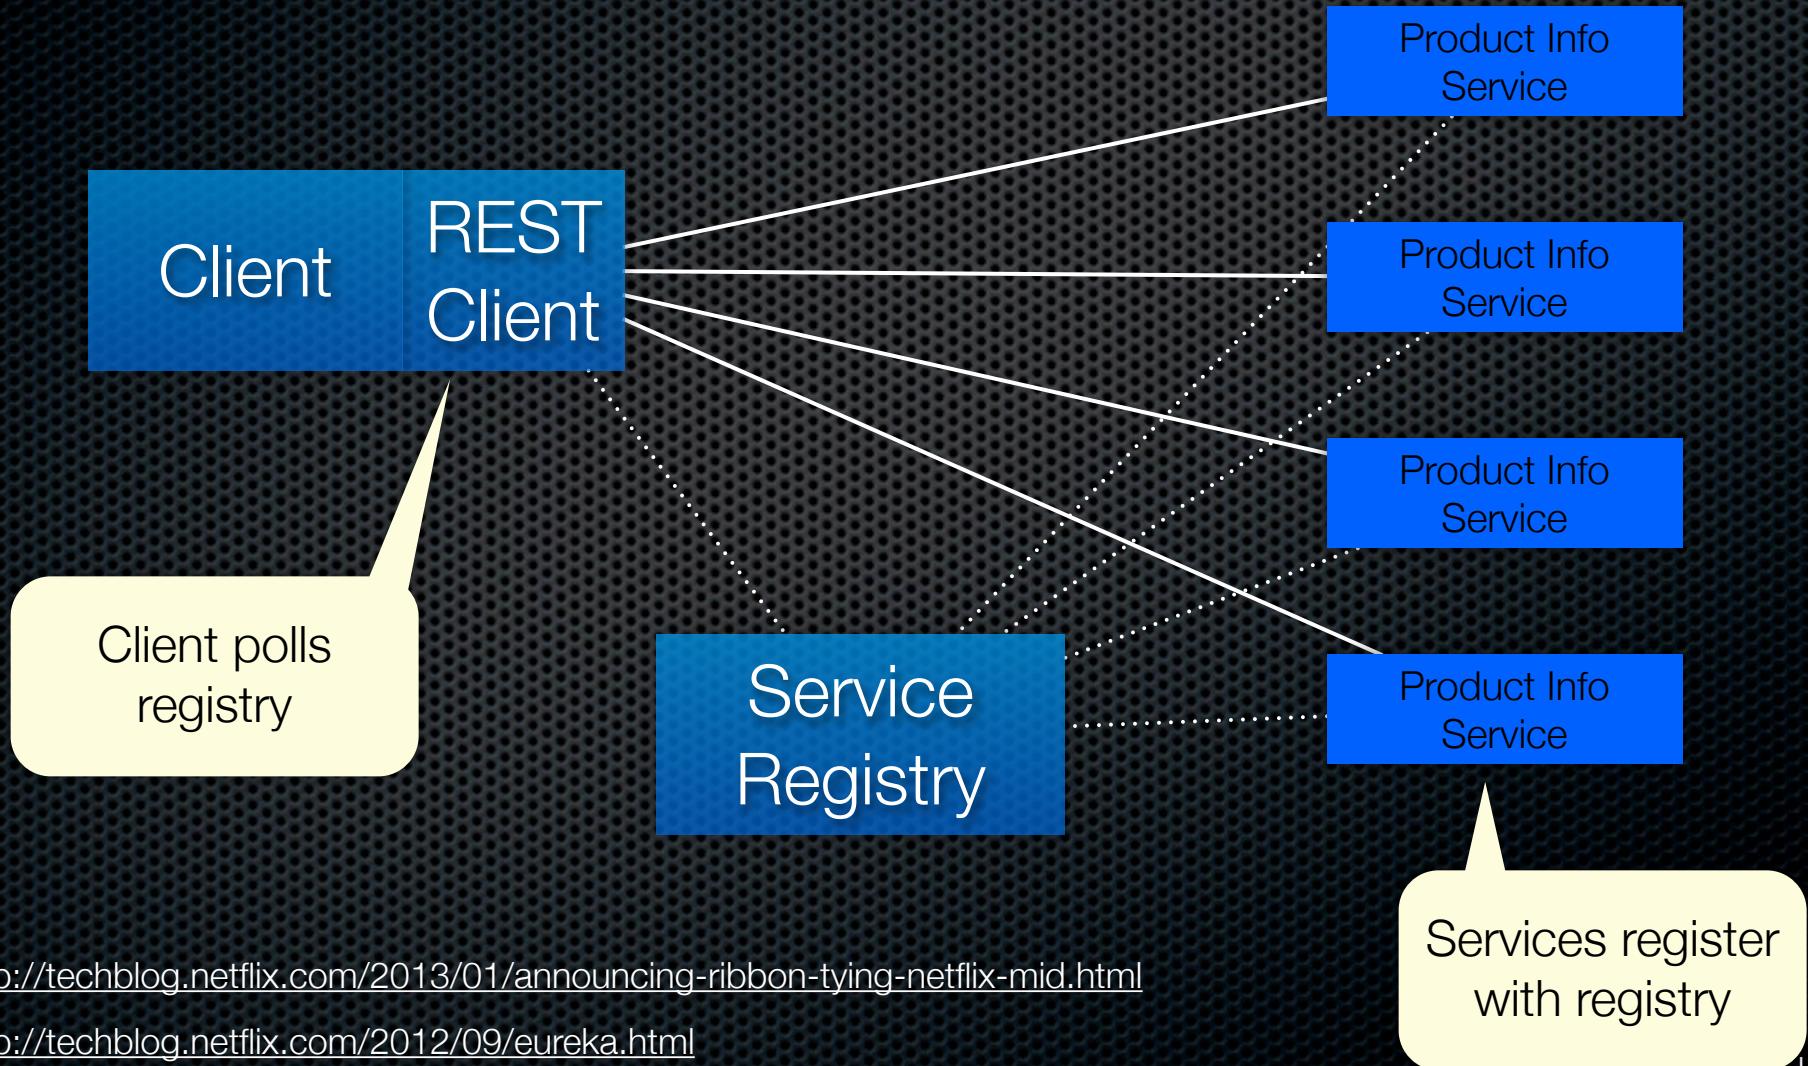
Pros and cons of HTTP

- ❖ Pros
 - ❖ Simple and familiar
 - ❖ Request/reply is easy
 - ❖ Firewall friendly
 - ❖ No intermediate broker
- ❖ Cons
 - ❖ Only supports request/reply
 - ❖ Server must be available
 - ❖ Client needs to discover URL(s) of server(s)

Discovery option #1: Internal load balancer



Discovery option #2: client-side load balancing

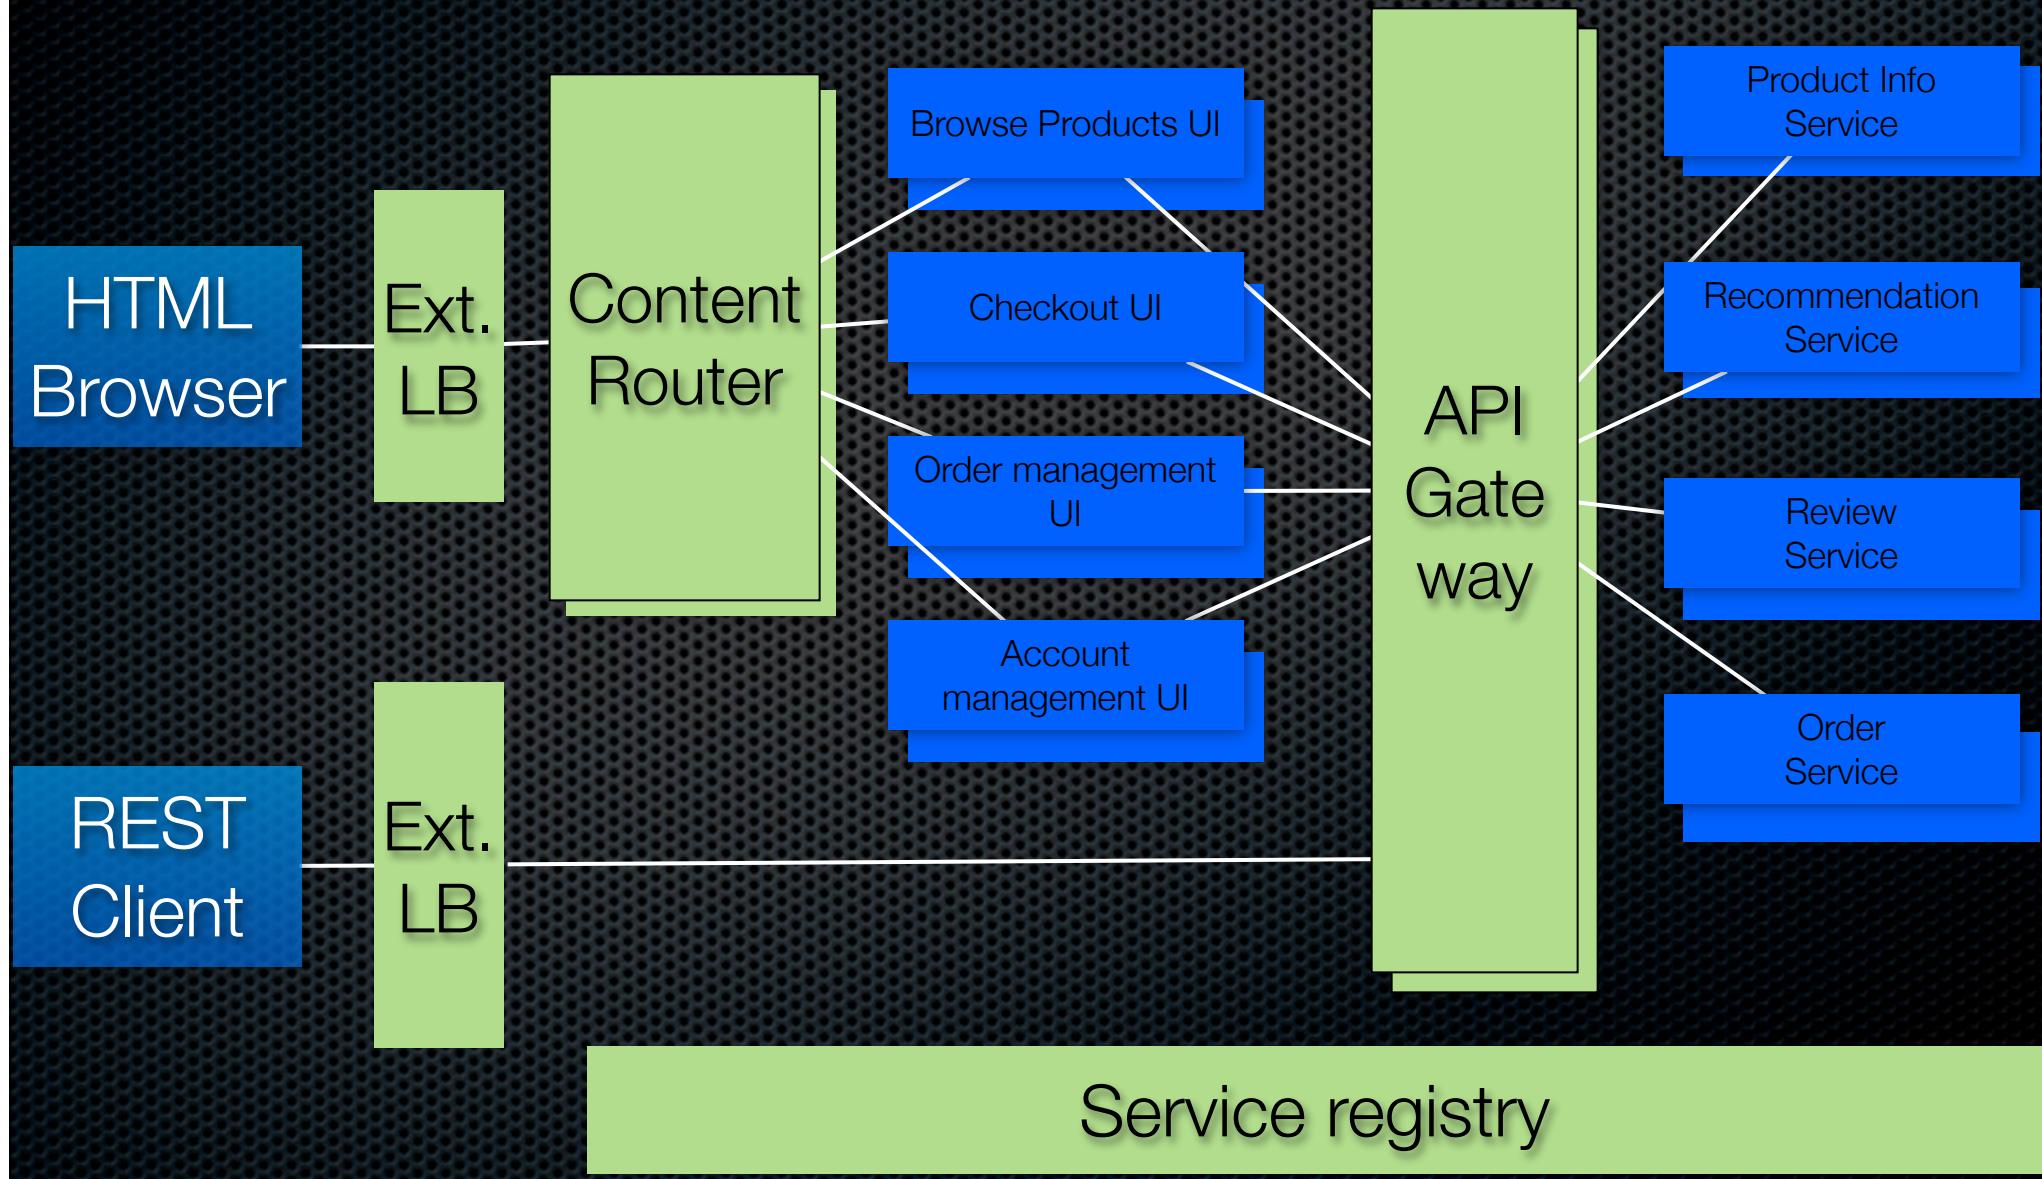


<http://techblog.netflix.com/2013/01/announcing-ribbon-tying-netflix-mid.html>

<http://techblog.netflix.com/2012/09/eureka.html>

@crichtson

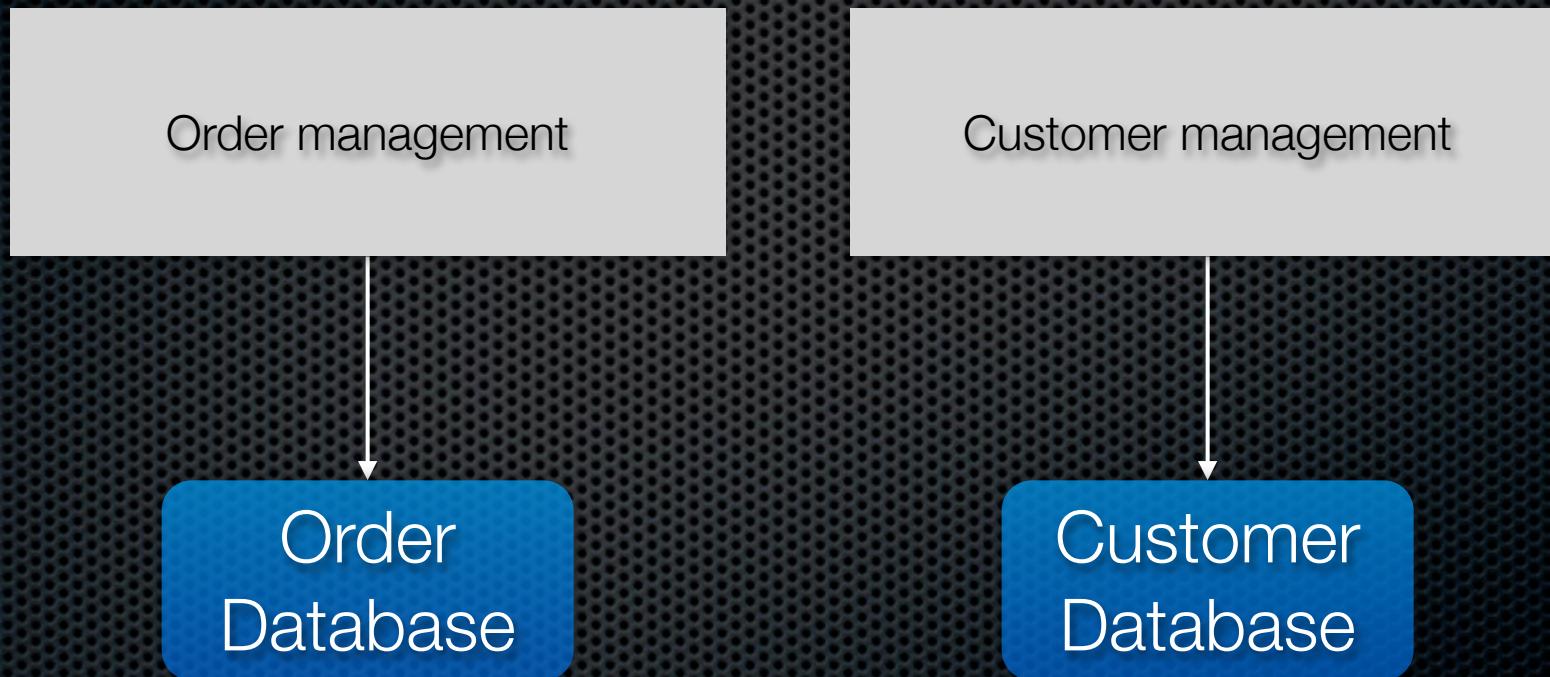
Lots of moving parts!



Agenda

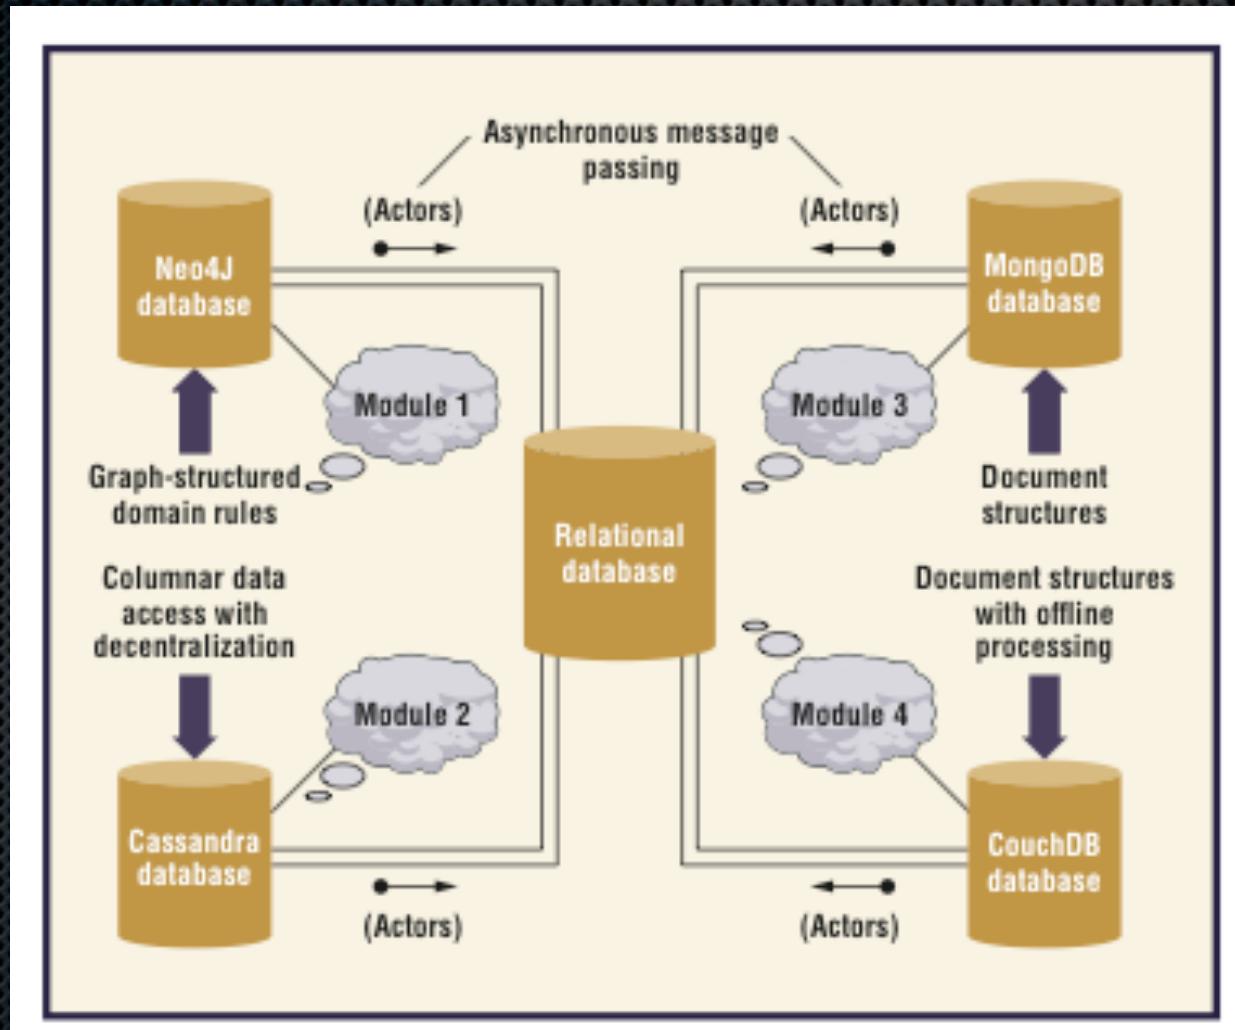
- The (sometimes evil) monolith
- Decomposing applications into services
- Client ⇔ service interaction design
- Decentralized data management

Decomposed services ⇒ decomposed databases



Separate databases ⇒ less coupling

Decomposed databases ⇒ polyglot persistence



Untangling orders and customers



Problems

- Reads
 - Service A needs to read data owned by service B
- Updates
 - Transaction must update data owned by multiple services

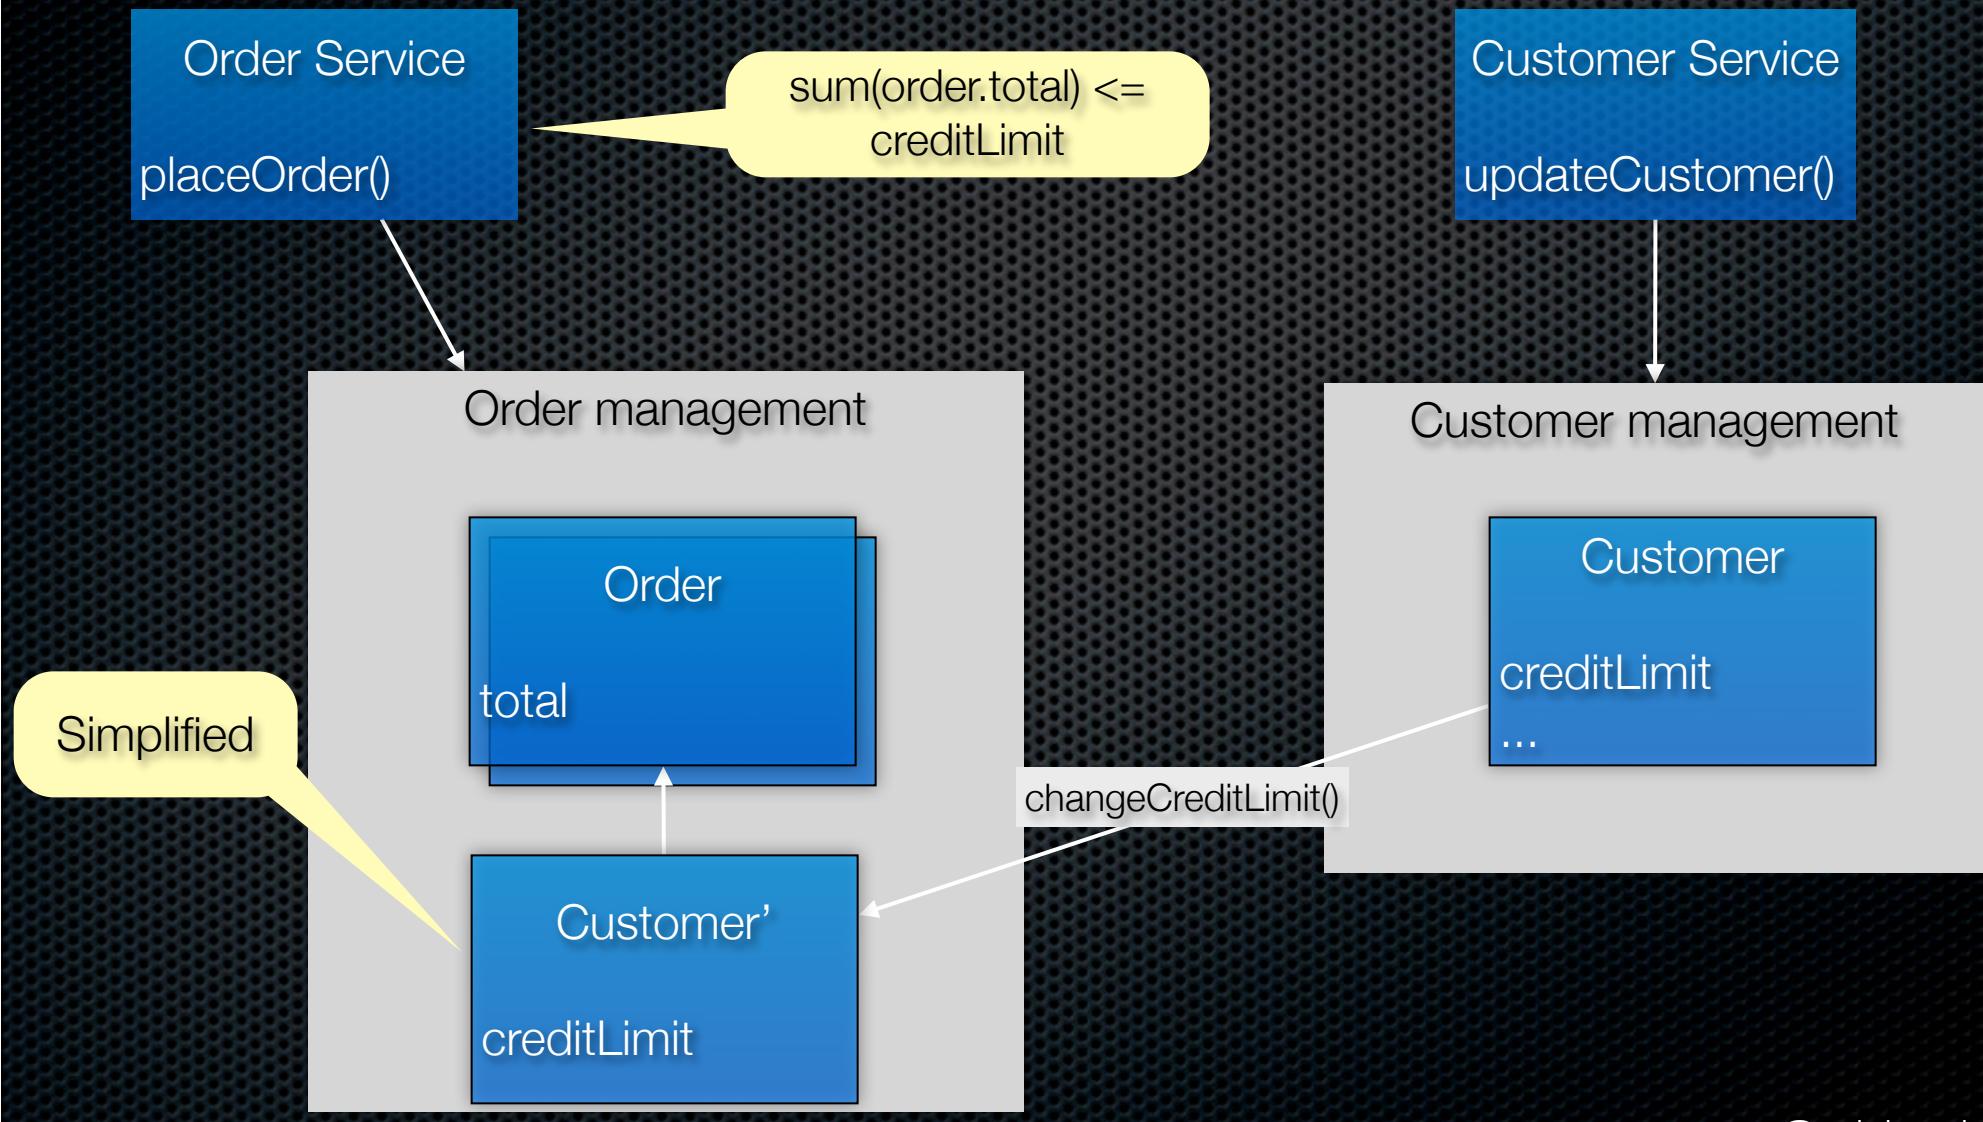
Handling reads: requesting credit limit



Pulling data

- ❖ Benefits
 - ❖ Simple to implement
 - ❖ Ensures data is fresh
- ❖ Drawbacks
 - ❖ Reduces availability
 - ❖ Increases response time

Handling reads: replicating the credit limit

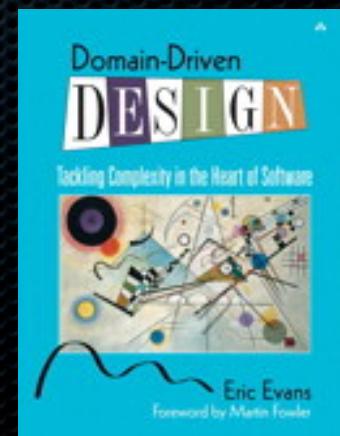


Useful idea: Bounded context

- Different services have a different view of a domain object, e.g.
 - User Management = complex view of user
 - Rest of application: User = PK + ACL + Name



- Different services can have a different domain model



Replicating data

- ❖ Benefits
 - ❖ Improved availability for reads
 - ❖ Improves latency
- ❖ Drawbacks
 - ❖ Additional complexity of replication mechanism

How to handle updates (including of replicated data)?

Use distributed transactions

- ❖ Benefits
 - ❖ Guarantees consistency
- ❖ Drawbacks
 - ❖ Complex
 - ❖ Reduced availability

Use eventual consistency

- ▣ How
 - ▣ Services publish events when data changes
 - ▣ Subscribing services update their data
- ▣ Benefits:
 - ▣ Simpler
 - ▣ Better availability
- ▣ Drawbacks:
 - ▣ Application has to handle inconsistencies

How do services publish events?

To maintain consistency the
application must
atomically publish an event
whenever
a domain object changes

Change tracking options

- ❖ Database triggers
- ❖ Hibernate event listener
- ❖ Ad hoc event publishing code mixed into business logic
- ❖ Domain events - “formal” modeling of events
- ❖ Event Sourcing

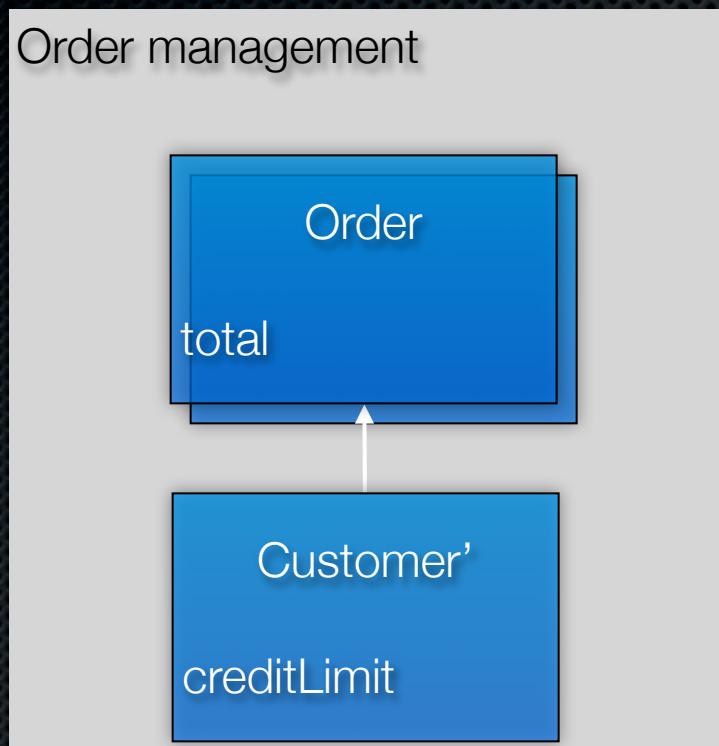
Event sourcing

- ❖ An event-centric approach to designing domain models
 - ❖ Aggregates handle commands by generating events
 - ❖ Apply events update aggregate state
- ❖ Persist events **NOT** state
 - ❖ Replay events to recreate the current state of an aggregate
- ❖ Event Store \approx database + message broker

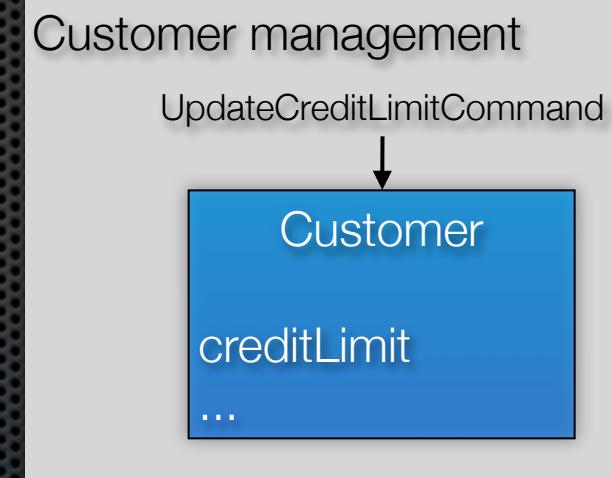
EventStore API

```
trait EventStore {  
  
    def save[T] (entityId: Id, events: Seq[Event]): T  
  
    def update[T] (entityId: Id,  
                  version: EntityVersion, events: Seq[Event]): T  
  
    def load[T] (entityType: Class[T], entityId: EntityId): T  
  
    def subscribe(...) : ...  
    ...  
}
```

Using event sourcing



Customer Service
updateCustomer()



CustomerCreditLimitUpdatedEvent

CustomerCreditLimitUpdatedEvent(...)

Event Store

Customer aggregate

```
case class Customer(customerId: String, creditLimit: BigDecimal)
  extends ValidatingAggregate[Customer, CustomerCommands.CustomerCommand] {

  def this() = this(null, null)

  override def validate = {
    case CreateCustomerCommand(customerId, creditLimit) =>
      Seq(CustomerCreatedEvent(customerId, creditLimit))
    case UpdateCreditLimitCommand(newLimit) if newLimit >= 0 =>
      Seq(CustomerCreditLimitUpdatedEvent(newLimit))
  }

  override def apply = {
    case CustomerCreatedEvent(customerId, creditLimit) =>
      copy(customerId=customerId, creditLimit=creditLimit)
    case CustomerCreditLimitUpdatedEvent(newLimit) =>
      copy(creditLimit=newLimit)
  }
}
```

Command
⇒
Events

Event
⇒
Updated state

Unfamiliar but it solves many problems

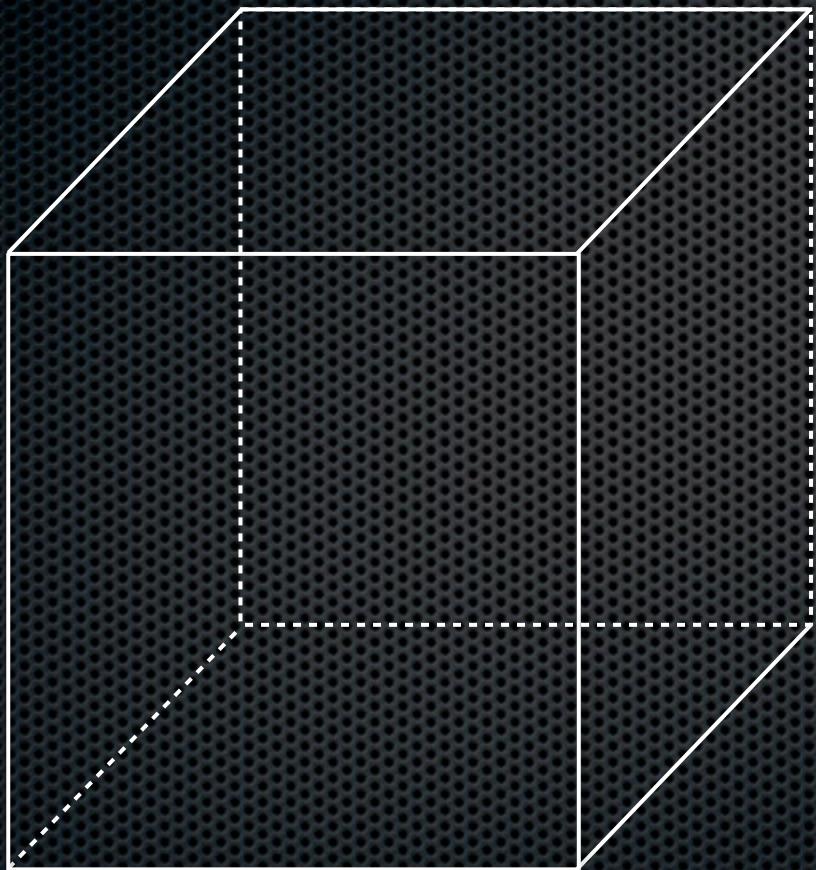
- ❖ Eliminates O/R mapping problem
- ❖ Supports both SQL and NoSQL databases
- ❖ Publishes events reliably
- ❖ Reliable eventual consistency framework
- ❖ ...

Summary

Monolithic applications are simple to develop and deploy

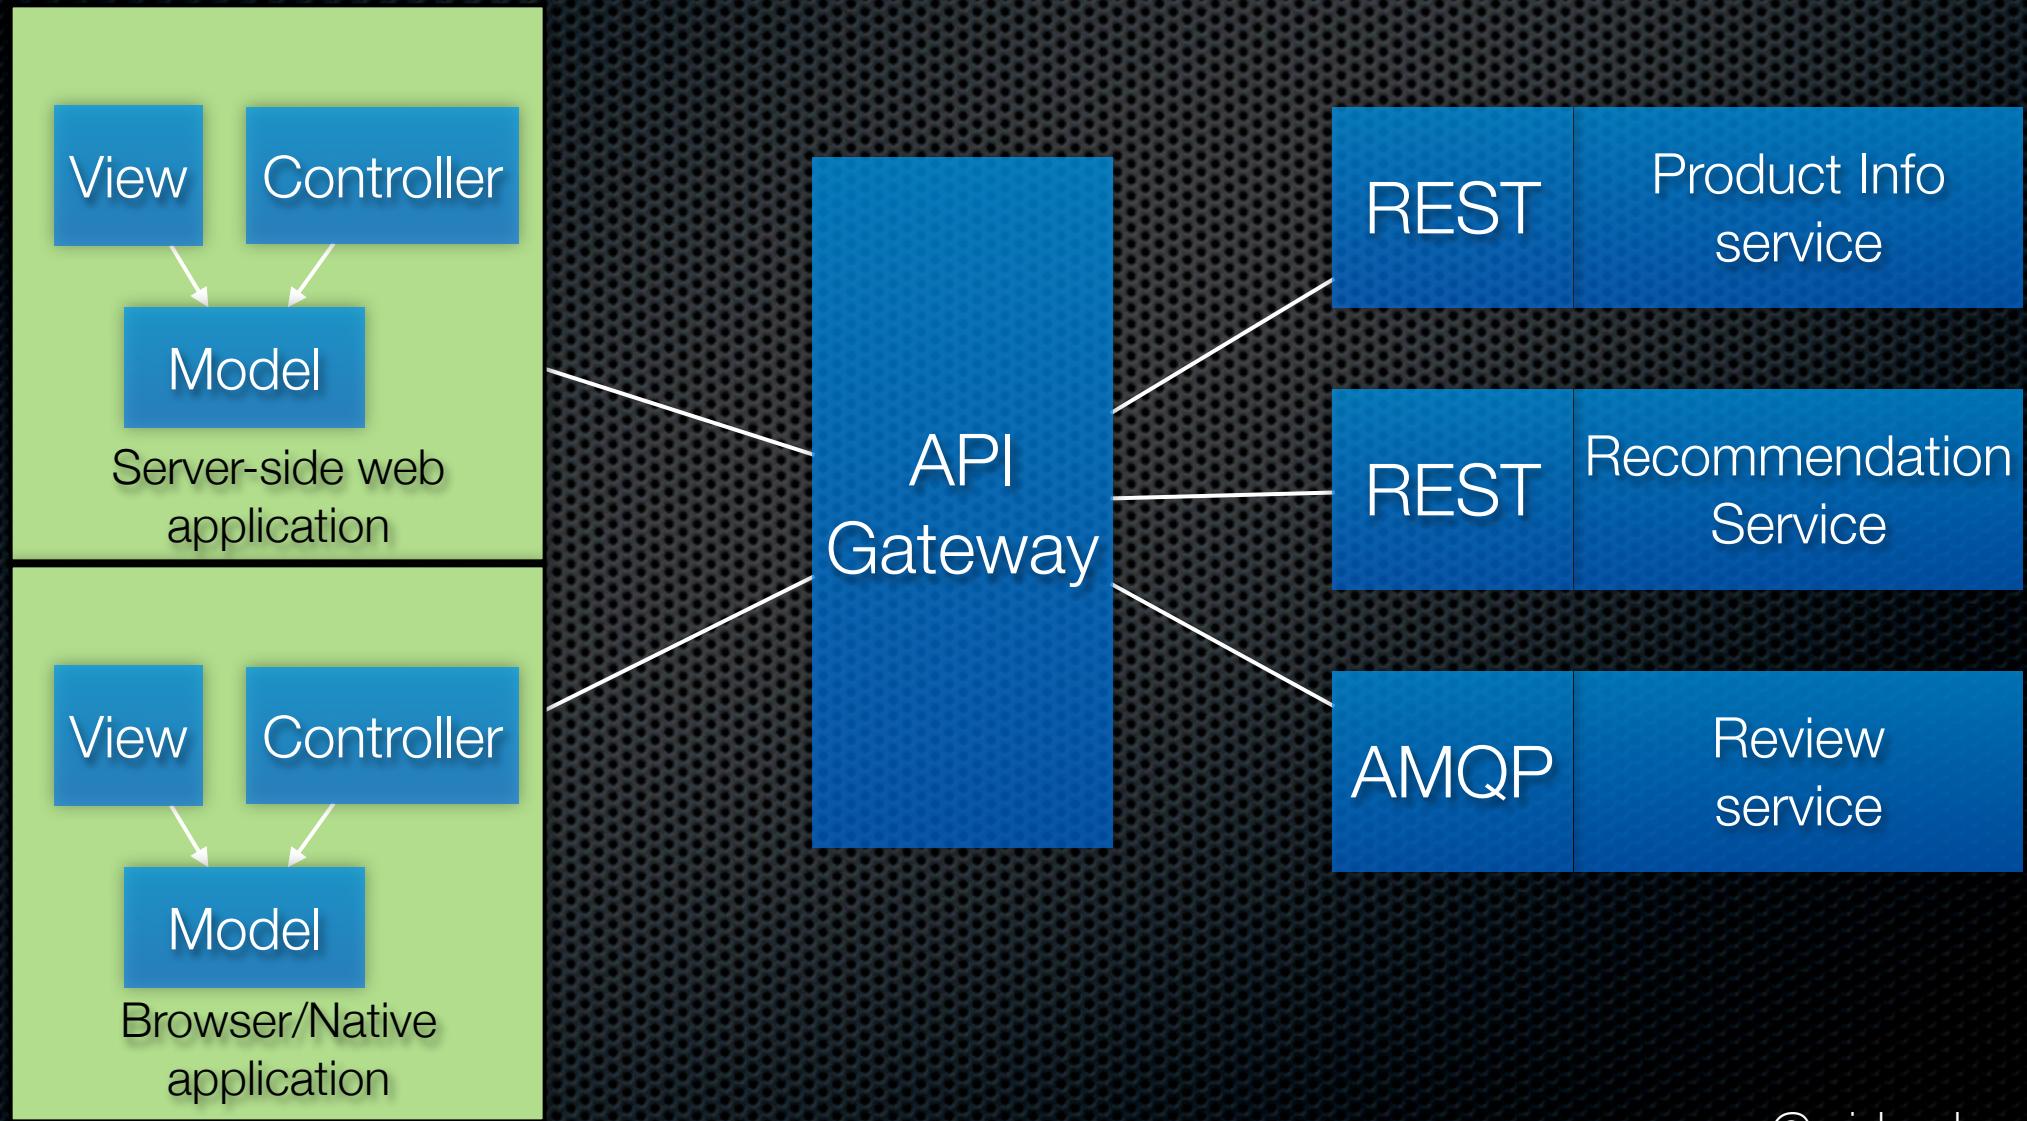
BUT have significant drawbacks

Apply the scale cube

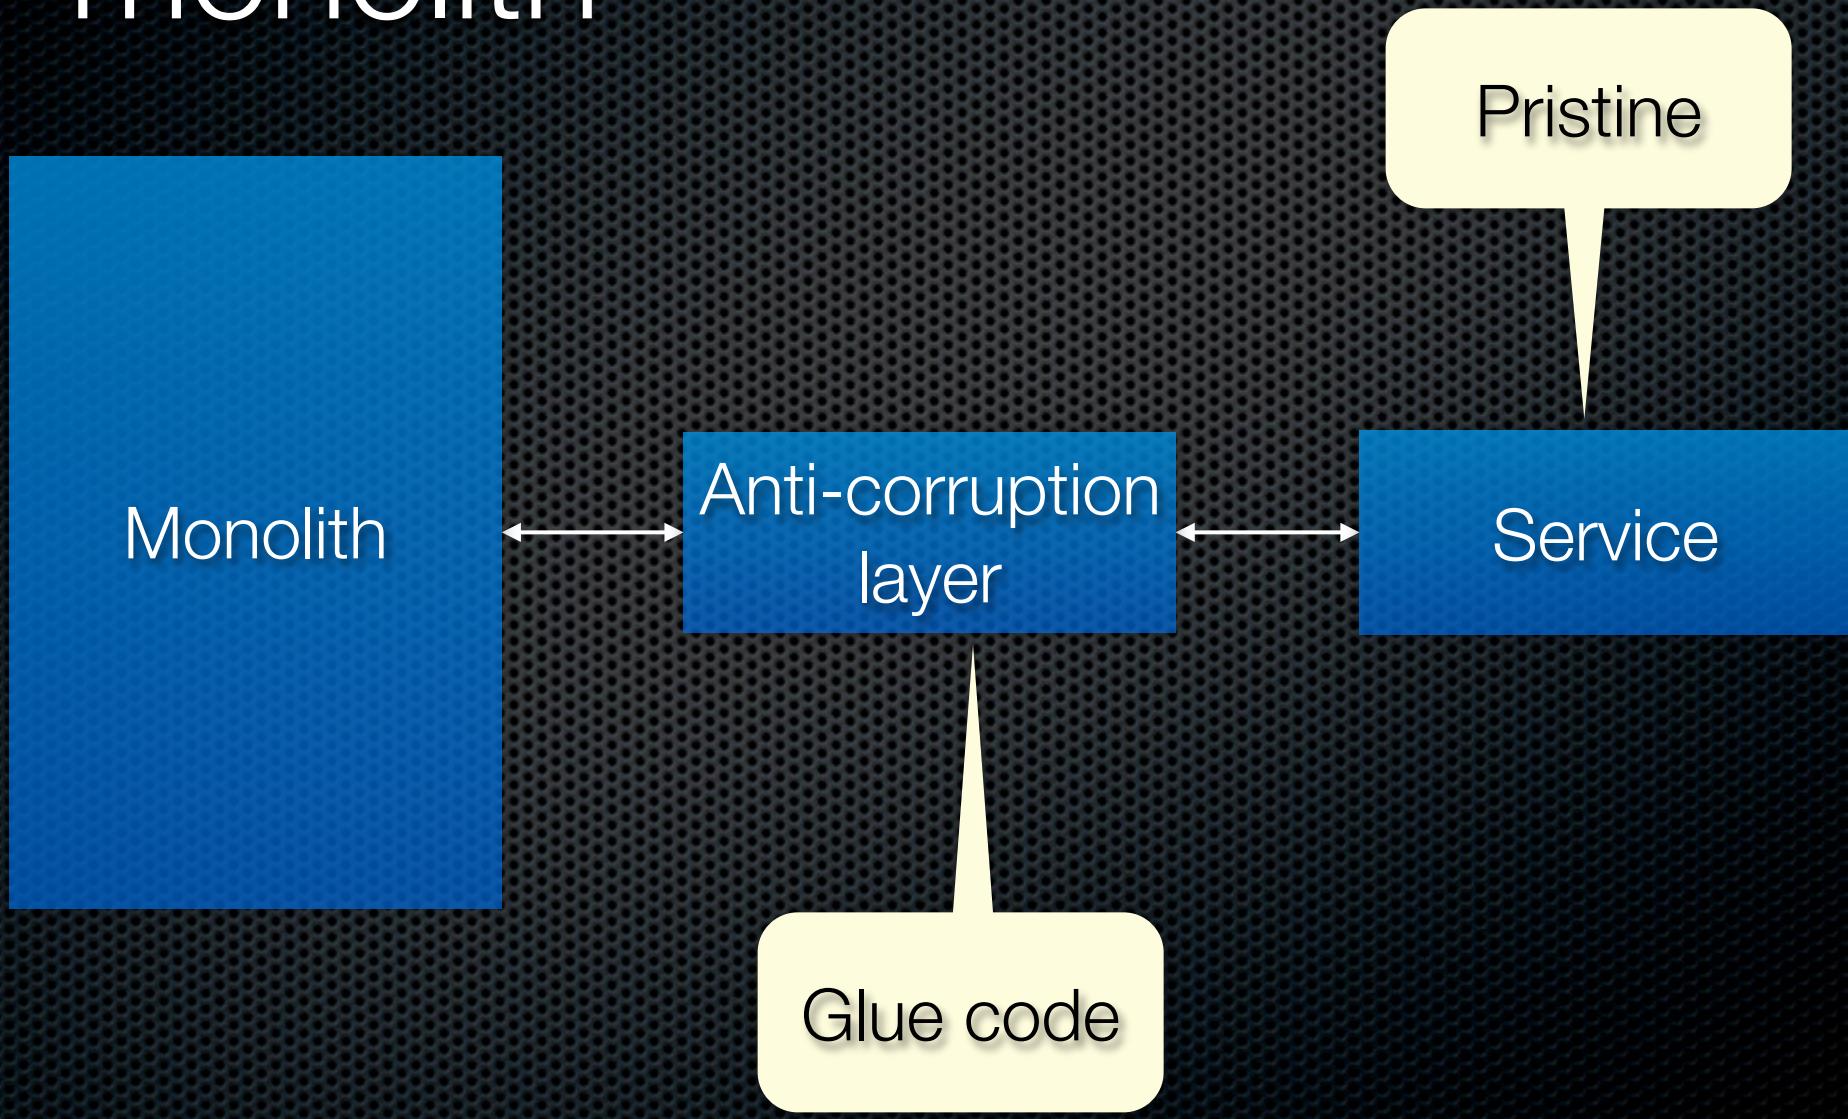


- Modular, polyglot, and scalable applications
- Services developed, deployed and scaled independently

Use a modular, polyglot architecture



Start refactoring your monolith



 @crichton

crichton@crichton.net



Questions?

<http://plainoldobjects.com>